

task

September 20, 2021

[]:

1 Abstract

Business problem - banks faceses heigh loss due to defaulters , we have to predict person will be defaulter or not with probablity of him to be defaulter

ML foormulaion - this is binary classification task - in this task we have predict probablity of person to be defaulter or not - in this problem it is important to have less false posiitve rate becace it can costs heigh loss to the bank - matix for evaluation is auc matrix - interpritability is important to give beacuse of which features preson should be defaulter one

first cut approach - first we will perform basic EDA - Handling of missing data - Handle imbalance dataset - after data perparation we will add some new features - Do required preprocessing over data (numerical ,categorical and text features) - apply ml model - predict submission

```
[ ]: %matplotlib inline
      %pip install chart-studio
      import warnings
      warnings.filterwarnings("ignore")

      import sqlite3
      import pandas as pd
      import numpy as np
      import nltk
      import string
      import matplotlib.pyplot as plt
      import seaborn as sns
      from sklearn.feature_extraction.text import TfidfTransformer
      from sklearn.feature_extraction.text import TfidfVectorizer

      from sklearn.feature_extraction.text import CountVectorizer
      from sklearn.metrics import confusion_matrix
      from sklearn import metrics
      from sklearn.metrics import roc_curve, auc
      from nltk.stem.porter import PorterStemmer
      from sklearn.impute import SimpleImputer
```

```

import re
# Tutorial about Python regular expressions: https://pymotw.com/2/re/
import string
from nltk.corpus import stopwords
from nltk.stem import PorterStemmer
from nltk.stem.wordnet import WordNetLemmatizer

from gensim.models import Word2Vec
from gensim.models import KeyedVectors
import pickle

from tqdm import tqdm
import os

from chart_studio.plotly import plotly
import plotly.offline as offline
import plotly.graph_objs as go
offline.init_notebook_mode()
from collections import Counter

```

Collecting chart-studio

Downloading chart_studio-1.1.0-py3-none-any.whl (64 kB)

| 64 kB 1.9 MB/s

Requirement already satisfied: plotly in /usr/local/lib/python3.7/dist-packages (from chart-studio) (4.4.1)

Requirement already satisfied: six in /usr/local/lib/python3.7/dist-packages (from chart-studio) (1.15.0)

Requirement already satisfied: requests in /usr/local/lib/python3.7/dist-packages (from chart-studio) (2.23.0)

Requirement already satisfied: retrying>=1.3.3 in /usr/local/lib/python3.7/dist-packages (from chart-studio) (1.3.3)

Requirement already satisfied: idna<3,>=2.5 in /usr/local/lib/python3.7/dist-packages (from requests->chart-studio) (2.10)

Requirement already satisfied: urllib3!=1.25.0,!1.25.1,<1.26,>=1.21.1 in /usr/local/lib/python3.7/dist-packages (from requests->chart-studio) (1.24.3)

Requirement already satisfied: chardet<4,>=3.0.2 in /usr/local/lib/python3.7/dist-packages (from requests->chart-studio) (3.0.4)

Requirement already satisfied: certifi>=2017.4.17 in /usr/local/lib/python3.7/dist-packages (from requests->chart-studio) (2021.5.30)

Installing collected packages: chart-studio

Successfully installed chart-studio-1.1.0

```
[ ]: train_data = pd.read_csv('train_indessa.csv')
test_data = pd.read_csv('test_indessa.csv')
```

```
[ ]: train_data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 532428 entries, 0 to 532427
```

```
Data columns (total 45 columns):
```

#	Column	Non-Null Count	Dtype
0	member_id	532428 non-null	int64
1	loan_amnt	532428 non-null	int64
2	funded_amnt	532428 non-null	int64
3	funded_amnt_inv	532428 non-null	float64
4	term	532428 non-null	object
5	batch_enrolled	447279 non-null	object
6	int_rate	532428 non-null	float64
7	grade	532428 non-null	object
8	sub_grade	532428 non-null	object
9	emp_title	501595 non-null	object
10	emp_length	505537 non-null	object
11	home_ownership	532428 non-null	object
12	annual_inc	532425 non-null	float64
13	verification_status	532428 non-null	object
14	pymnt_plan	532428 non-null	object
15	desc	75599 non-null	object
16	purpose	532428 non-null	object
17	title	532338 non-null	object
18	zip_code	532428 non-null	object
19	addr_state	532428 non-null	object
20	dti	532428 non-null	float64
21	delinq_2yrs	532412 non-null	float64
22	inq_last_6mths	532412 non-null	float64
23	mths_since_last_delinq	259874 non-null	float64
24	mths_since_last_record	82123 non-null	float64
25	open_acc	532412 non-null	float64
26	pub_rec	532412 non-null	float64
27	revol_bal	532428 non-null	float64
28	revol_util	532141 non-null	float64
29	total_acc	532412 non-null	float64
30	initial_list_status	532428 non-null	object
31	total_rec_int	532428 non-null	float64
32	total_rec_late_fee	532428 non-null	float64
33	recoveries	532428 non-null	float64
34	collection_recovery_fee	532428 non-null	float64
35	collections_12_mths_ex_med	532333 non-null	float64
36	mths_since_last_major_derog	132980 non-null	float64
37	application_type	532428 non-null	object
38	verification_status_joint	305 non-null	object
39	last_week_pay	532428 non-null	object
40	acc_now_delinq	532412 non-null	float64
41	tot_coll_amt	490424 non-null	float64
42	tot_cur_bal	490424 non-null	float64

```

43 total_rev_hi_lim          490424 non-null float64
44 loan_status              532428 non-null int64
dtypes: float64(23), int64(4), object(18)
memory usage: 182.8+ MB

```

```

[ ]: dtypes = train_data.dtypes
categorical = dtypes[dtypes=='object']
print('following are categorical features')
catgs = np.array(categorical.keys())
cat_feature = catgs[catgs != 'desc']
print(cat_feature)

```

```

following are categorical features
['term' 'batch_enrolled' 'grade' 'sub_grade' 'emp_title' 'emp_length'
 'home_ownership' 'verification_status' 'pymnt_plan' 'purpose' 'title'
 'zip_code' 'addr_state' 'initial_list_status' 'application_type'
 'verification_status_joint' 'last_week_pay']

```

```

[ ]: print('desc is text feature')

```

```

desc is text feature

```

```

[ ]: dtypes = train_data.dtypes
categorical = dtypes[dtypes!='object']
print('following are numeric features')
int_features = np.array(categorical.keys())

print(int_features , int_features.shape[0])

```

```

following are numeric features
['member_id' 'loan_amnt' 'funded_amnt' 'funded_amnt_inv' 'int_rate'
 'annual_inc' 'dti' 'delinq_2yrs' 'inq_last_6mths'
 'mths_since_last_delinq' 'mths_since_last_record' 'open_acc' 'pub_rec'
 'revol_bal' 'revol_util' 'total_acc' 'total_rec_int' 'total_rec_late_fee'
 'recoveries' 'collection_recovery_fee' 'collections_12_mths_ex_med'
 'mths_since_last_major_derog' 'acc_now_delinq' 'tot_coll_amt'
 'tot_cur_bal' 'total_rev_hi_lim' 'loan_status'] 27

```

```

[ ]: print("Number of data points in train data", train_data.shape)
print('-'*50)
print("The attributes of data :", train_data.columns.values)

```

```

Number of data points in train data (532428, 45)
-----

```

```

The attributes of data : ['member_id' 'loan_amnt' 'funded_amnt'
 'funded_amnt_inv' 'term'
 'batch_enrolled' 'int_rate' 'grade' 'sub_grade' 'emp_title' 'emp_length'
 'home_ownership' 'annual_inc' 'verification_status' 'pymnt_plan' 'desc'

```

```
'purpose' 'title' 'zip_code' 'addr_state' 'dti' 'delinq_2yrs'
'inq_last_6mths' 'mths_since_last_delinq' 'mths_since_last_record'
'open_acc' 'pub_rec' 'revol_bal' 'revol_util' 'total_acc'
'initial_list_status' 'total_rec_int' 'total_rec_late_fee' 'recoveries'
'collection_recovery_fee' 'collections_12_mths_ex_med'
'mths_since_last_major_derog' 'application_type'
'verification_status_joint' 'last_week_pay' 'acc_now_delinq'
'tot_coll_amt' 'tot_cur_bal' 'total_rev_hi_lim' 'loan_status']
```

```
[ ]: y_value_counts = train_data['loan_status'].value_counts()
print("Number of members thar are Defaulters ", y_value_counts[1], ", (", "\u2192(y_value_counts[1]/(y_value_counts[1]+y_value_counts[0]))*100,"%")")
print("Number of projects thar are non defaulters ", y_value_counts[0], ", (", "\u2192(y_value_counts[0]/(y_value_counts[1]+y_value_counts[0]))*100,"%")")

fig, ax = plt.subplots(figsize=(6, 6), subplot_kw=dict(aspect="equal"))
recipe = ["Defaulters", "Non Defaulters"]

data = [y_value_counts[1], y_value_counts[0]]

wedges, texts = ax.pie(data, wedgeprops=dict(width=0.5), startangle=-40)

bbox_props = dict(boxstyle="square,pad=0.3", fc="w", ec="k", lw=0.72)
kw = dict(xycoords='data', textcoords='data', arrowprops=dict(arrowstyle="-"),
          bbox=bbox_props, zorder=0, va="center")

for i, p in enumerate(wedges):
    ang = (p.theta2 - p.theta1)/2. + p.theta1
    y = np.sin(np.deg2rad(ang))
    x = np.cos(np.deg2rad(ang))
    horizontalalignment = {-1: "right", 1: "left"}[int(np.sign(x))]
    connectionstyle = "angle,angleA=0,angleB={}".format(ang)
    kw["arrowprops"].update({"connectionstyle": connectionstyle})
    ax.annotate(recipe[i], xy=(x, y), xytext=(1.35*np.sign(x), 1.4*y),
                horizontalalignment=horizontalalignment, **kw)

ax.set_title("All defaulter and non defaulters")

plt.show()
```

```
Number of members thar are Defaulters 125827 , ( 23.632678972555915 %)
Number of projects thar are non defaulters 406601 , ( 76.36732102744409 %)
```

All defaulter and non defaulters



1.0.1 we can say that data is imbalance data

1. we will choose to undersample to deal with an imbalance dataset
2. we will choose over sampling to deal with imbalance data because with undersampling we don't want to lose the information
3. we will go to do oversampling of minority class upto 2:3 ratio

```
[ ]: blank_spaces = train_data.eq('').sum()  
blank_spaces
```

```
[ ]: member_id          0  
loan_amnt              0  
funded_amnt           0  
funded_amnt_inv       0  
term                  0  
batch_enrolled      106079  
int_rate              0  
grade                 0  
sub_grade             0  
emp_title             0  
emp_length            0
```

home_ownership	0
annual_inc	0
verification_status	0
pymnt_plan	0
desc	0
purpose	0
title	0
zip_code	0
addr_state	0
dti	0
delinq_2yrs	0
inq_last_6mths	0
mths_since_last_delinq	0
mths_since_last_record	0
open_acc	0
pub_rec	0
revol_bal	0
revol_util	0
total_acc	0
initial_list_status	0
total_rec_int	0
total_rec_late_fee	0
recoveries	0
collection_recovery_fee	0
collections_12_mths_ex_med	0
mths_since_last_major_derog	0
application_type	0
verification_status_joint	0
last_week_pay	0
acc_now_delinq	0
tot_coll_amt	0
tot_cur_bal	0
total_rev_hi_lim	0
loan_status	0
dtype: int64	

```
[ ]: # we will check nan values in overall dataset
      # review the columns with heighest amount of nan values
      na_values = train_data.isna().sum()
      na_values
```

[]: member_id	0
loan_amnt	0
funded_amnt	0
funded_amnt_inv	0
term	0
batch_enrolled	85149

int_rate	0
grade	0
sub_grade	0
emp_title	30833
emp_length	26891
home_ownership	0
annual_inc	3
verification_status	0
pymnt_plan	0
desc	456829
purpose	0
title	90
zip_code	0
addr_state	0
dti	0
delinq_2yrs	16
inq_last_6mths	16
mths_since_last_delinq	272554
mths_since_last_record	450305
open_acc	16
pub_rec	16
revol_bal	0
revol_util	287
total_acc	16
initial_list_status	0
total_rec_int	0
total_rec_late_fee	0
recoveries	0
collection_recovery_fee	0
collections_12_mths_ex_med	95
mths_since_last_major_derog	399448
application_type	0
verification_status_joint	532123
last_week_pay	0
acc_now_delinq	16
tot_coll_amt	42004
tot_cur_bal	42004
total_rev_hi_lim	42004
loan_status	0
dtype:	int64

```
[ ]: #columns with nan value
train_data.isna().any(axis=0).sum().sum()
```

```
[ ]: 21
```

- so we can find there are huge no of nan values in some columns so we will going to exclude

those columns whose contains more that 50 % of nan values

- for features whose are beetwen 0 to 50 % of nan values we will going replace them with median if they are numerical features and if they are categorical then we will going to replace them by most occuring category

```
[ ]: # Creating a dictionary whose keys are the column names and values are the
      ↪percentage of missing values
x = train_data
nan_count = {k: list(x.isna().sum()*100/x.shape[0])[i] for i,k in enumerate(x.
      ↪columns)}


# Sorting the dictionary in descending order based on the percentage of missing
      ↪values
nan_count = {k: v for k, v in sorted(nan_count.items(), key=lambda item:
      ↪item[1],reverse=True)}

# Plotting a graph showing the top 15 features having highest percentage of
      ↪missing values
sns.set_style(style="whitegrid")
plt.figure(figsize=(20,10))

# Bar Plot
plot = sns.barplot(x= list(nan_count.keys())[:15],y = list(nan_count.values())[:
      ↪15],palette="hls")

# Add annotations above each bar signifying their value
for p in plot.patches:
    plot.annotate('{:.1f}%'.format(p.get_height()), (p.get_x()+0.2, p.
      ↪get_height()+1))

# Make y-axis more interpretable
plot.set_yticklabels(map('{:.1f}%'.format, plot.yaxis.get_majorticklocs()))
plt.show()
```



```
[ ]: print('columns whose having nan values more than 50% ')
x = nan_count
sorted_x = sorted(x.items(), key=lambda kv: kv[1], reverse=True)
heigh_na_features = {i:j for i,j in sorted_x if j > 50}
print(heigh_na_features)
print('columns whose having nan values less than 50%')
medium_na_features = {i:j for i,j in sorted_x if j <= 50 and j > 0}
print(medium_na_features)
```

```
columns whose having nan values more than 50%
{'verification_status_joint': 99.94271525915241, 'desc': 85.80108484151847,
'mths_since_last_record': 84.57575484384743, 'mths_since_last_major_derog':
75.02385299045129, 'mths_since_last_delinq': 51.19077133434004}
columns whose having nan values less than 50%
{'batch_enrolled': 15.99258491288963, 'tot_coll_amt': 7.889141818236457,
'tot_cur_bal': 7.889141818236457, 'total_rev_hi_lim': 7.889141818236457,
'emp_title': 5.791017752635098, 'emp_length': 5.05063595453282, 'revol_util':
0.05390400204346879, 'collections_12_mths_ex_med': 0.017842788132855524,
'title': 0.01690369402059997, 'delinq_2yrs': 0.0030051011592177723,
'inq_last_6mths': 0.0030051011592177723, 'open_acc': 0.0030051011592177723,
'pub_rec': 0.0030051011592177723, 'total_acc': 0.0030051011592177723,
'acc_now_delinq': 0.0030051011592177723, 'annual_inc': 0.0005634564673533323}
```

```
[ ]: train_data_preprocessed = train_data
test_data_preprocessed = test_data
```

```
[ ]: #drop the columns having missing values rate > 50%
```

```
train_data_preprocessed = train_data_preprocessed.
↳drop(columns=['verification_status_joint', 'desc', 'mths_since_last_record',
↳'mths_since_last_major_derog', 'mths_since_last_delinq'])
```

```
[ ]: train_data_preprocessed.shape
```

```
[ ]: (532428, 40)
```

```
[ ]:
```

2 Handle missing values

- batch batch_enrolled feature contains nan values + blank spaces
- this is total nearly 35.9 % of total rows
- as we are not domain expert we cant tell values whcih are missing are missing reandomly or the missing value have its own meaning
- there are three options in this case
 - remove column
 - replace spaces and nan with most frequest value
 - conside blank space and nan as new categories
- we will choose third option

```
[ ]: (blank_spaces['batch_enrolled']+na_values['batch_enrolled'])/train_data.shape[0]
```

```
[ ]: 0.3591621777968101
```

```
[ ]: train_data_preprocessed['batch_enrolled'] =
↳train_data_preprocessed['batch_enrolled'].fillna('unknown')
```

```
[ ]: train_data_preprocessed['batch_enrolled'] =
↳train_data_preprocessed['batch_enrolled'].map(lambda x: 'blank' if x == ' '
↳else x)
```

```
[ ]: batch_enrolled = train_data_preprocessed['batch_enrolled']
print(batch_enrolled[batch_enrolled=='blank'].shape ,
↳batch_enrolled[batch_enrolled=='unknown'].shape)
```

```
(106079,) (85149,)
```

```
[ ]: train_data_preprocessed['batch_enrolled']
```

```
[ ]: 0          blank
1      BAT1586599
2      BAT1586599
3      BAT4808022
4      BAT2833642
```

```

...
532423      blank
532424      BAT2003848
532425      unknown
532426      BAT3193689
532427      BAT4136152
Name: batch_enrolled, Length: 532428, dtype: object

```

```

[ ]: medium_na_features
categorical_na_features = list(filter(lambda x: x in
    ↳ cat_feature, list(medium_na_features.keys()))
int_na_features = list(filter(lambda x: x in
    ↳ int_features, list(medium_na_features.keys()))

```

```

[ ]: categorical_na_features

```

```

[ ]: ['batch_enrolled', 'emp_title', 'emp_length', 'title']

```

```

[ ]: int_na_features

```

```

[ ]: ['tot_coll_amt',
      'tot_cur_bal',
      'total_rev_hi_lim',
      'revol_util',
      'collections_12_mths_ex_med',
      'delinq_2yrs',
      'inq_last_6mths',
      'open_acc',
      'pub_rec',
      'total_acc',
      'acc_now_delinq',
      'annual_inc']

```

```

[ ]: def replace_with_most_frequent (feature):
      values_data = train_data_preprocessed[feature].value_counts()
      max_value = values_data[values_data == values_data.max()].index[0]
      return train_data_preprocessed[feature].fillna(max_value[0]) , max_value

```

```

[ ]: def fill_cat_na(categorical_na_features):
      save_feature_max_value = dict({})
      for i in categorical_na_features:
          feature , val = replace_with_most_frequent(i)
          train_data_preprocessed[i] = feature
          save_feature_max_value[i] = val
      pickle.dump(save_feature_max_value, file = open("save_feature_max_value.
    ↳ pickle", "wb"))

```

```

save_feature_max_value = pickle.load(open("save_feature_max_value.pickle",
↪"rb"))
return save_feature_max_value

```

```
[ ]: fill_cat_na(categorical_na_features)
```

```
[ ]: {'batch_enrolled': 'blank',
      'emp_length': '10+ years',
      'emp_title': 'Teacher',
      'title': 'Debt consolidation'}
```

```
[ ]: save_feature_max_value = pickle.load(open("save_feature_max_value.pickle",
↪"rb"))
save_feature_max_value

```

```
[ ]: {'batch_enrolled': 'blank',
      'emp_length': '10+ years',
      'emp_title': 'Teacher',
      'title': 'Debt consolidation'}
```

```
[ ]: def replace_with_median (feature):
      median_value = train_data_preprocessed[feature].median()
      train_feature = train_data_preprocessed[feature].fillna(median_value)
      return train_feature , median_value

```

```
[ ]: def fill_int_na(int_na_features):
      save_feature_median_value = dict({})
      for i in int_na_features:
          feature , val = replace_with_median(i)
          train_data_preprocessed[i] = feature
          save_feature_median_value[i] = val
      pickle.dump(save_feature_median_value, file = open("save_feature_median_value.
↪pickle", "wb"))
      save_feature_median_value = pickle.load(open("save_feature_median_value.
↪pickle", "rb"))
      return save_feature_median_value

```

```
[ ]: fill_int_na(int_na_features)
```

```
[ ]: {'acc_now_delinq': 0.0,
      'annual_inc': 65000.0,
      'collections_12_mths_ex_med': 0.0,
      'delinq_2yrs': 0.0,
      'inq_last_6mths': 0.0,
      'open_acc': 11.0,
      'pub_rec': 0.0,
      'revol_util': 56.0,

```

```
'tot_coll_amt': 0.0,
'tot_cur_bal': 80669.5,
'total_acc': 24.0,
'total_rev_hi_lim': 23700.0}
```

```
[ ]: train_data_preprocessed.isna().sum()
```

```
[ ]: member_id          0
    loan_amnt           0
    funded_amnt         0
    funded_amnt_inv     0
    term                0
    batch_enrolled      0
    int_rate            0
    grade               0
    sub_grade           0
    emp_title           0
    emp_length          0
    home_ownership      0
    annual_inc          0
    verification_status 0
    pymnt_plan          0
    purpose             0
    title               0
    zip_code            0
    addr_state          0
    dti                 0
    delinq_2yrs         0
    inq_last_6mths      0
    open_acc            0
    pub_rec             0
    revol_bal           0
    revol_util          0
    total_acc           0
    initial_list_status 0
    total_rec_int       0
    total_rec_late_fee  0
    recoveries          0
    collection_recovery_fee 0
    collections_12_mths_ex_med 0
    application_type    0
    last_week_pay       0
    acc_now_delinq      0
    tot_coll_amt        0
    tot_cur_bal         0
    total_rev_hi_lim    0
    loan_status         0
```

```
dtype: int64
```

```
#Basic EDA
```

```
[ ]: dtypes = train_data_preprocessed.dtypes
categorical = dtypes[dtypes=='object']
print('following are categorical features')
catgs = np.array(categorical.keys())
preprocess_cat_feature = catgs[catgs != 'desc']
print(preprocess_cat_feature ,preprocess_cat_feature.shape)
```

following are categorical features

```
['term' 'batch_enrolled' 'grade' 'sub_grade' 'emp_title' 'emp_length'
 'home_ownership' 'verification_status' 'pymnt_plan' 'purpose' 'title'
 'zip_code' 'addr_state' 'initial_list_status' 'application_type'
 'last_week_pay'] (16,)
```

```
[ ]: dtypes = train_data_preprocessed.dtypes
categorical = dtypes[dtypes!='object']
print('following are numerical features')
catgs = np.array(categorical.keys())
preprocess_int_feature = catgs[np.logical_and(catgs != 'member_id', catgs !=_
↳'loan_status')]
print(preprocess_int_feature.shape)
preprocess_int_feature
```

following are numerical features
(22,)

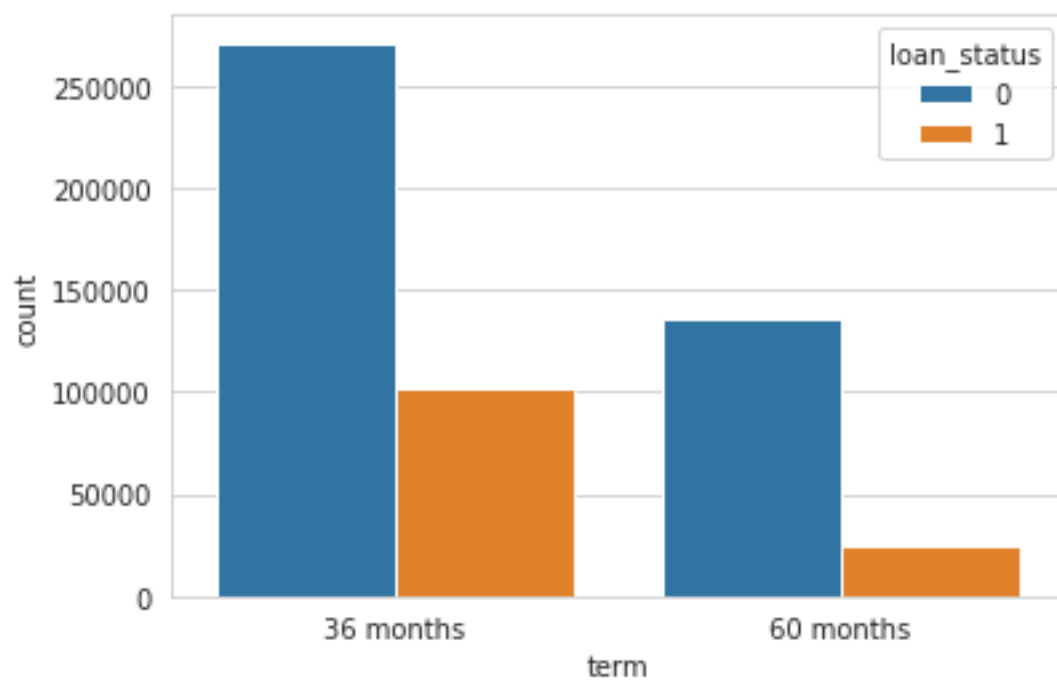
```
[ ]: array(['loan_amnt', 'funded_amnt', 'funded_amnt_inv', 'int_rate',
 'annual_inc', 'dti', 'delinq_2yrs', 'inq_last_6mths', 'open_acc',
 'pub_rec', 'revol_bal', 'revol_util', 'total_acc', 'total_rec_int',
 'total_rec_late_fee', 'recoveries', 'collection_recovery_fee',
 'collections_12_mths_ex_med', 'acc_now_delinq', 'tot_coll_amt',
 'tot_cur_bal', 'total_rev_hi_lim'], dtype=object)
```

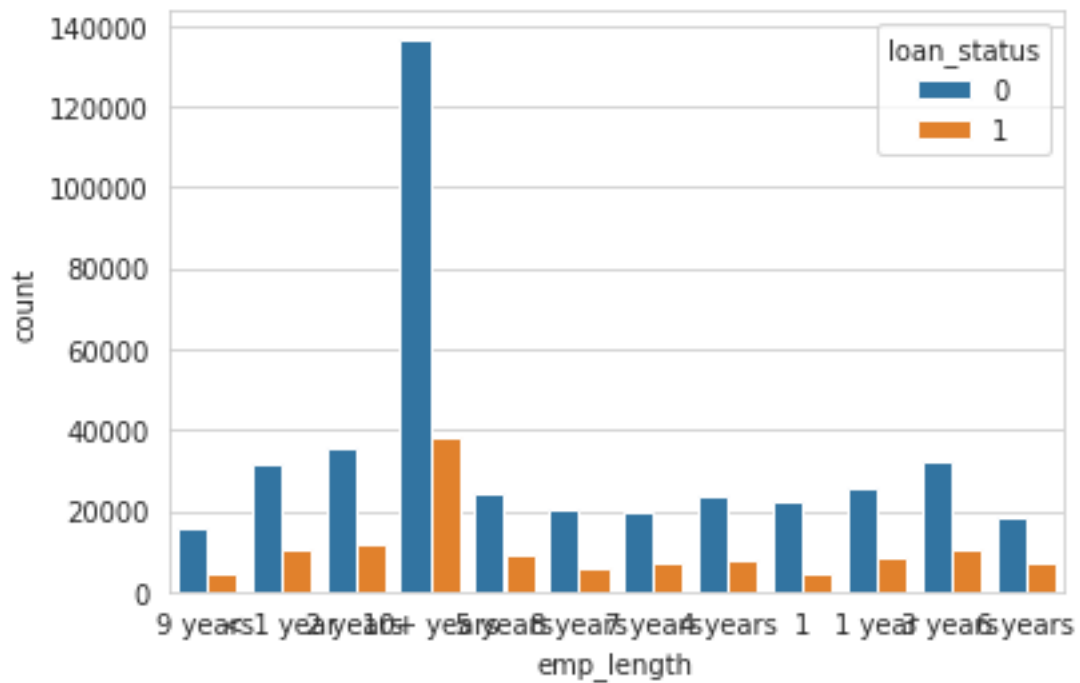
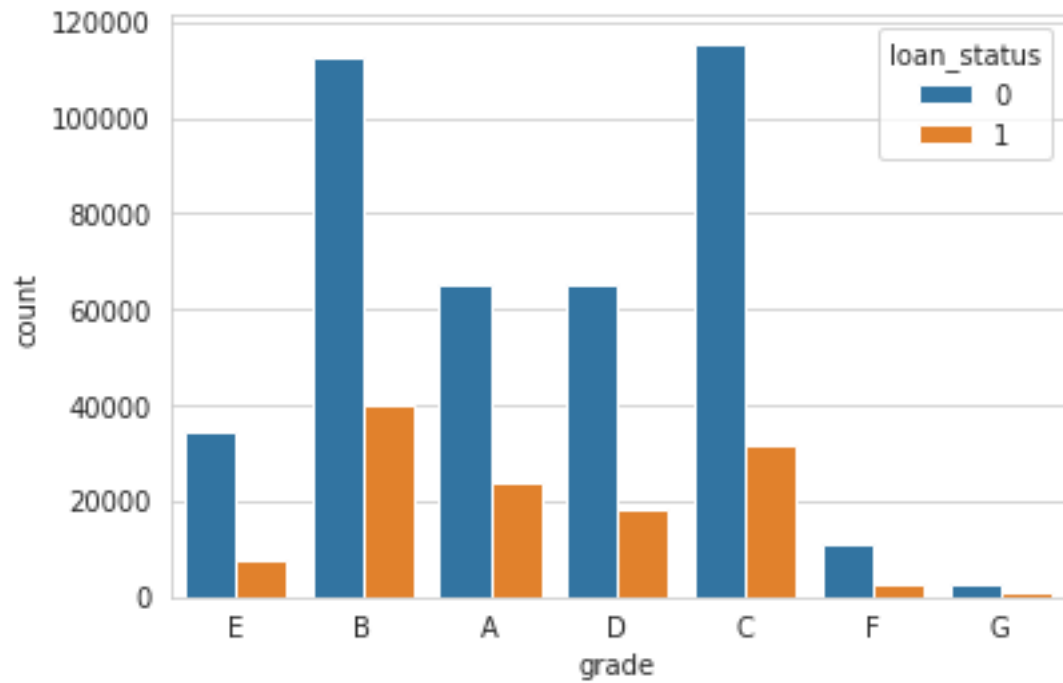
```
[ ]: def printCounterPlot(feature , row):
    # count plot on two categorical variable
    if train_data_preprocessed[feature].value_counts().shape[0] > 30:
        return
    sns.countplot(x = feature, hue = "loan_status", data =_
↳train_data_preprocessed)
    plt.show()
```

```
[ ]: train_data_preprocessed.loan_status.value_counts()
```

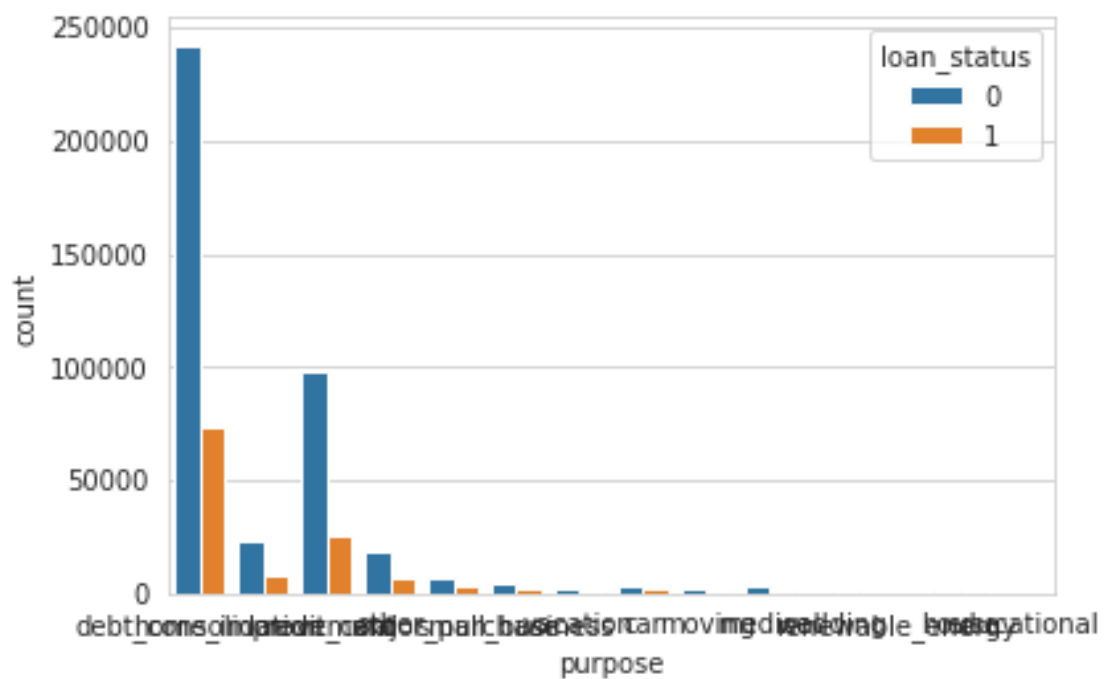
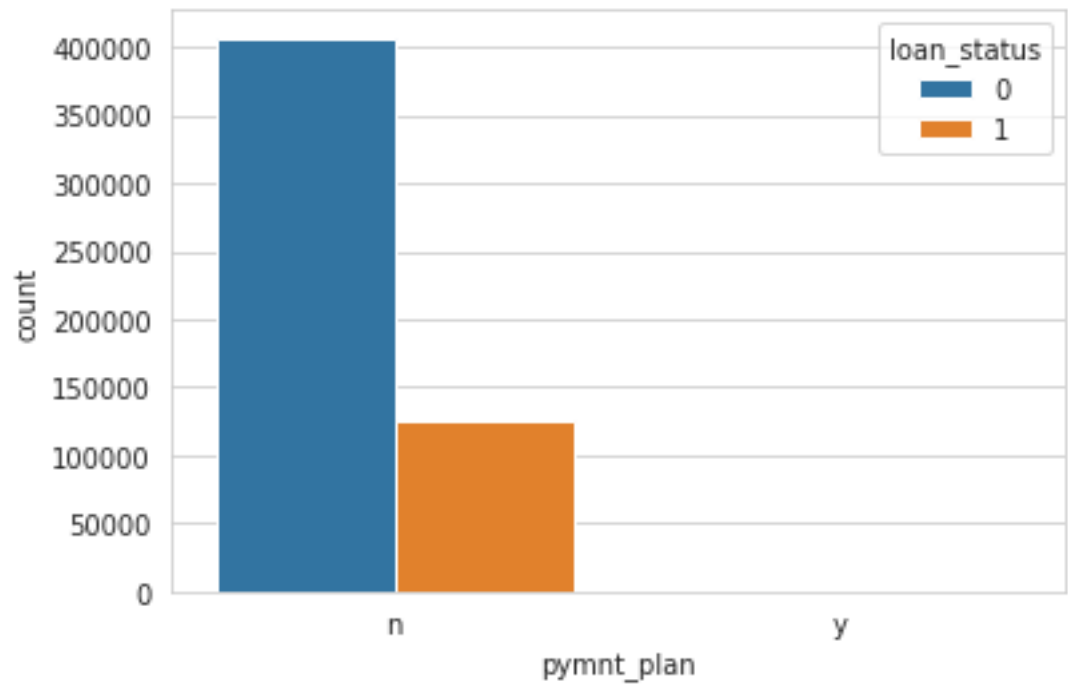
```
[ ]: 0    406601  
     1    125827  
     Name: loan_status, dtype: int64
```



```
[ ]: ### Data Visualization libraries  
import seaborn as sns  
import matplotlib.pyplot as plt  
  
for idx,cat_col in enumerate(preprocess_cat_feature):  
    printCounterPlot(cat_col,idx)  
#plt.subplots_adjust(hspace=1)
```











```
[ ]: train_data_preprocessed.application_type.value_counts()
```

```
[ ]: INDIVIDUAL    532123
      JOINT        305
      Name: application_type, dtype: int64
```

```
[ ]: application_type = train_data_preprocessed[train_data_preprocessed.loan_status_
      ↪== 0].application_type
      application_type[application_type == 'JOINT'].shape[0]
```

```
[ ]: 304
```

```
[ ]: application_type = train_data_preprocessed[train_data_preprocessed.loan_status_
      ↪== 1].application_type
      application_type[application_type == 'JOINT'].shape[0]
```

```
[ ]: 1
```

```
[ ]: train_data_preprocessed.pymnt_plan.value_counts()
```

```
[ ]: n    532420
      y      8
      Name: pymnt_plan, dtype: int64
```

```
[ ]: pymnt_plan = train_data_preprocessed[train_data_preprocessed.loan_status == 0].
      ↪pymnt_plan
      application_type[application_type == 'y'].shape[0]
```

```
[ ]: 0
```

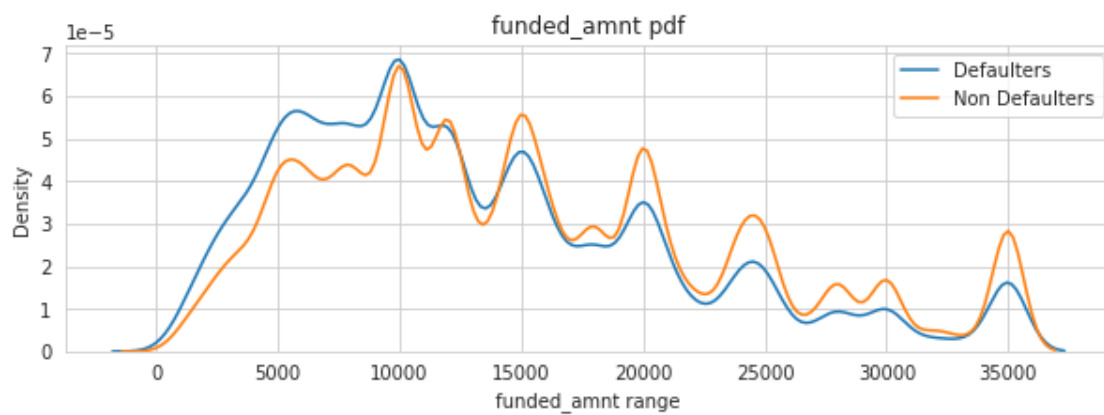
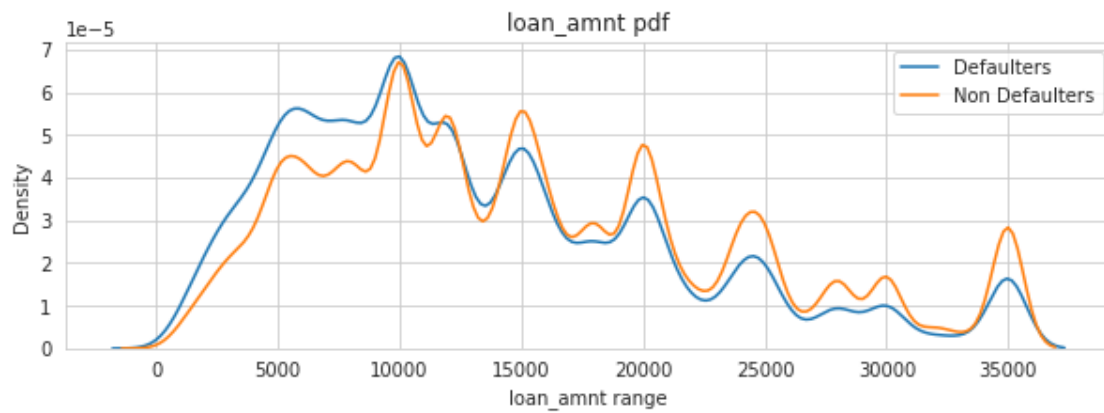
```
[ ]: train_data_preprocessed = train_data_preprocessed.
      ↪drop(columns=['application_type', 'pymnt_plan'])
```

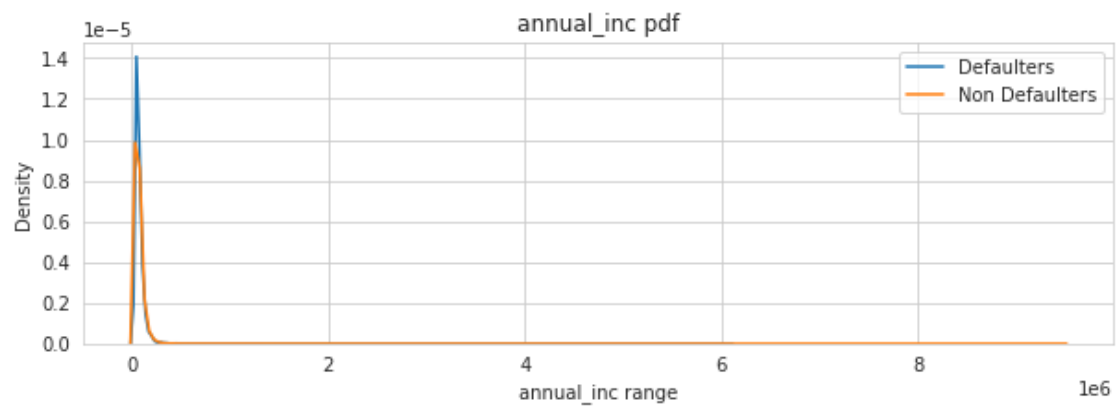
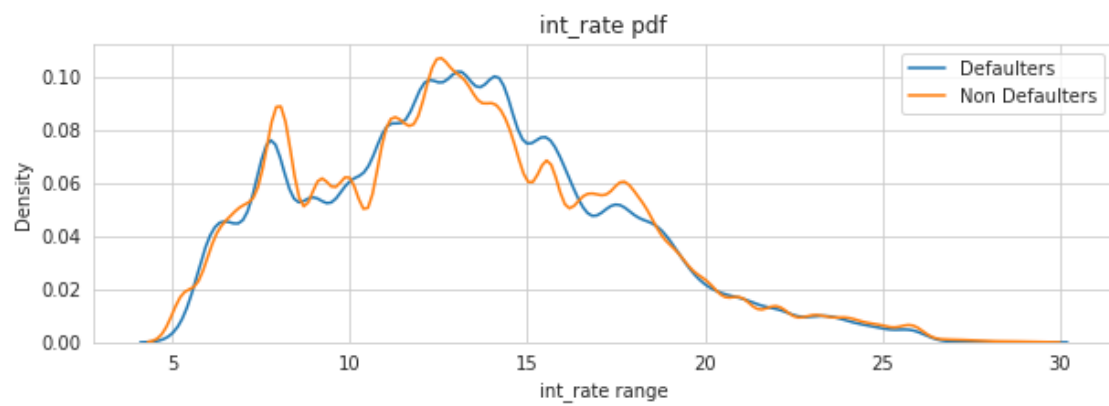

- in our dataset if we see mostly all members having individual application type we so if we drop this feature it will not affect result much
- same is about pymnt_plan feature almost all the members having pymnt_plan category n so this feature can misguide us
- we will drop this two features from categorical features

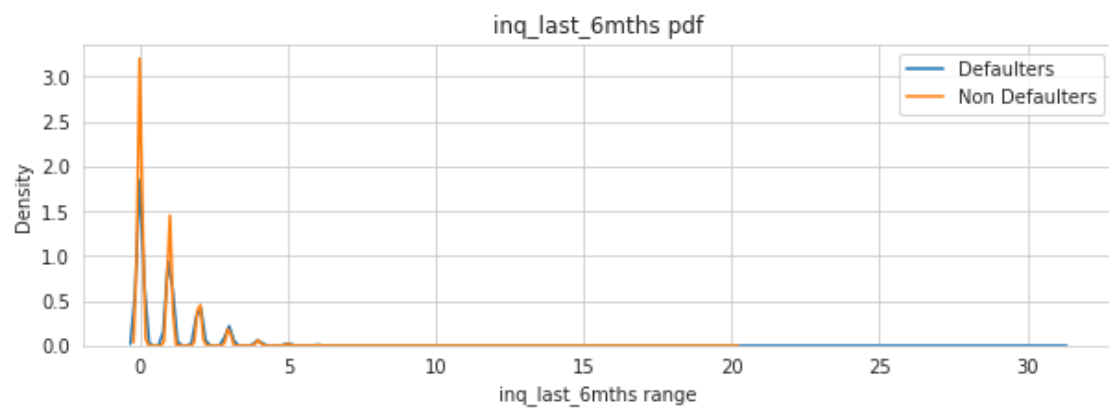
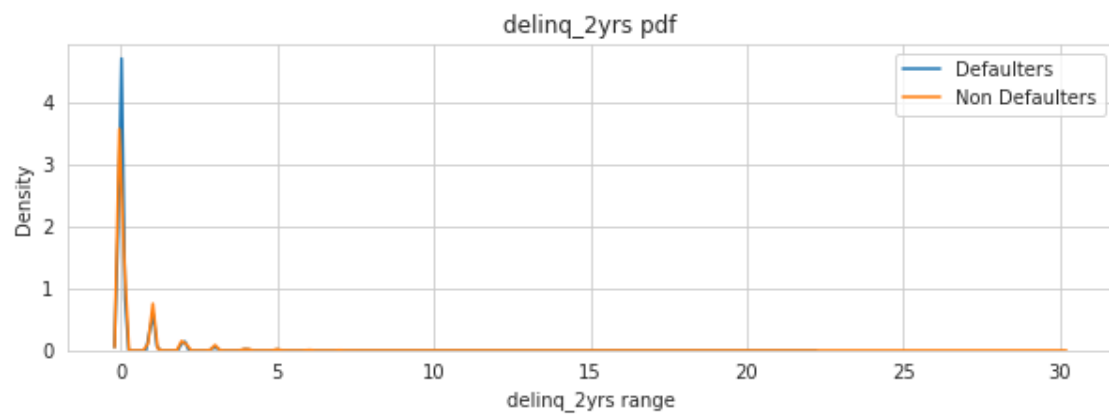
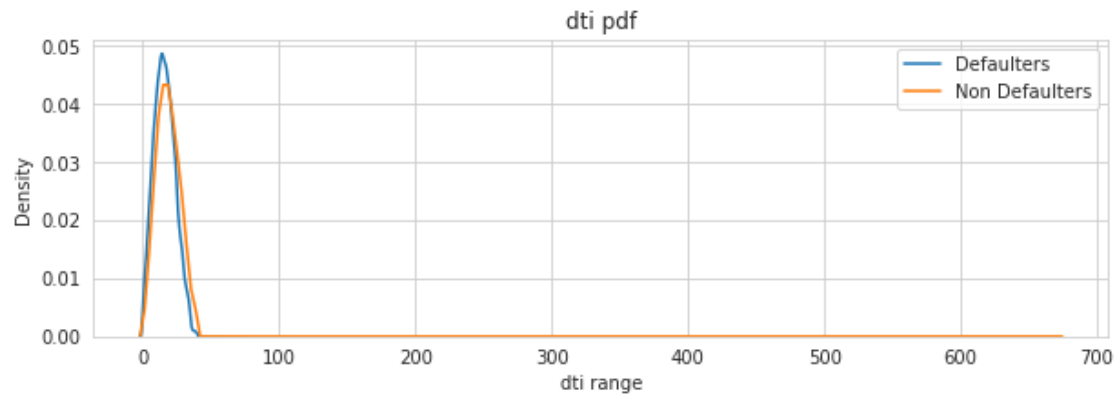
```
[ ]: def plotPdf(feature):
      defaulter = ↪
      ↪train_data_preprocessed[train_data_preprocessed['loan_status']==1][feature].
      ↪values
      non_defaulter = ↪
      ↪train_data_preprocessed[train_data_preprocessed['loan_status']==0][feature].
      ↪values
      plt.figure(figsize=(10,3))
      sns.distplot(defaulter, hist=False, label="Defaulters")
      sns.distplot(non_defaulter, hist=False, label="Non Defaulters")
```


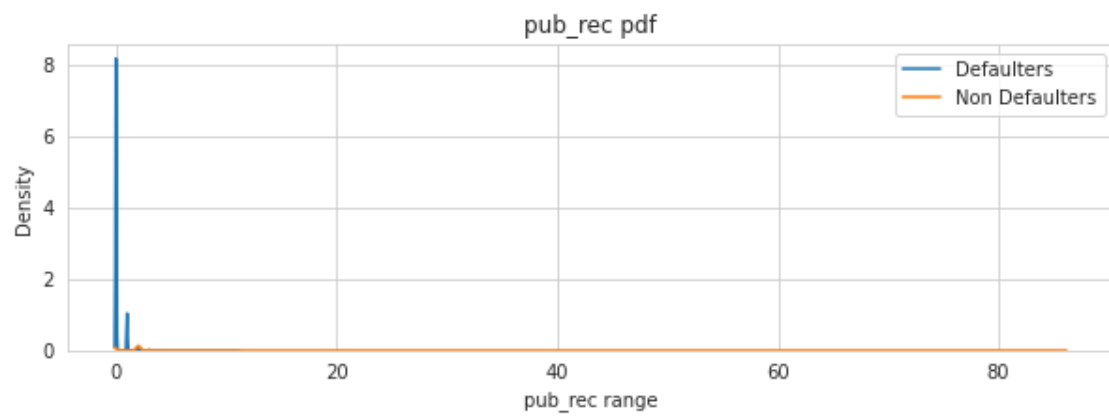
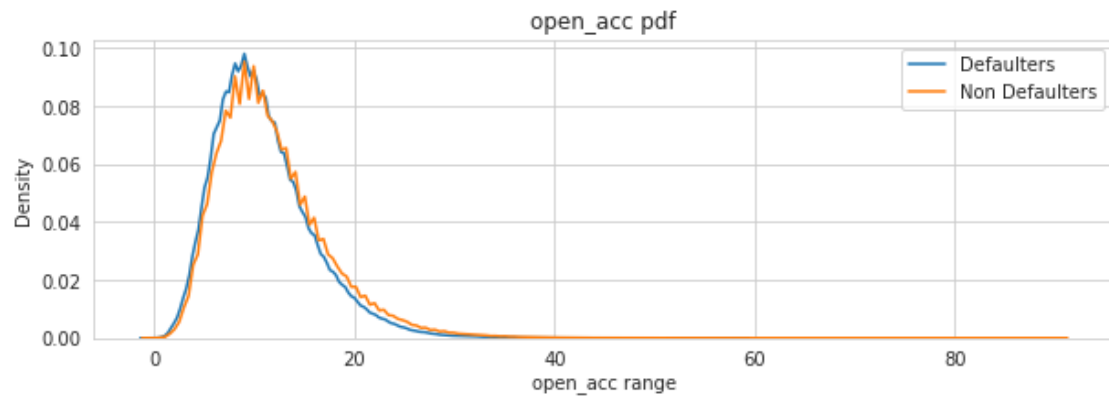
```
plt.title(feature+' pdf')
plt.xlabel(feature+' range')
plt.legend()
plt.show()
```

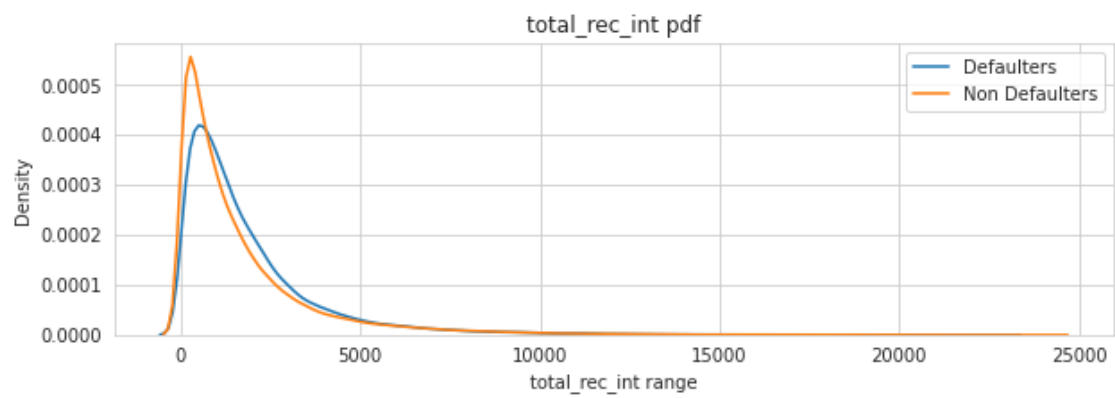
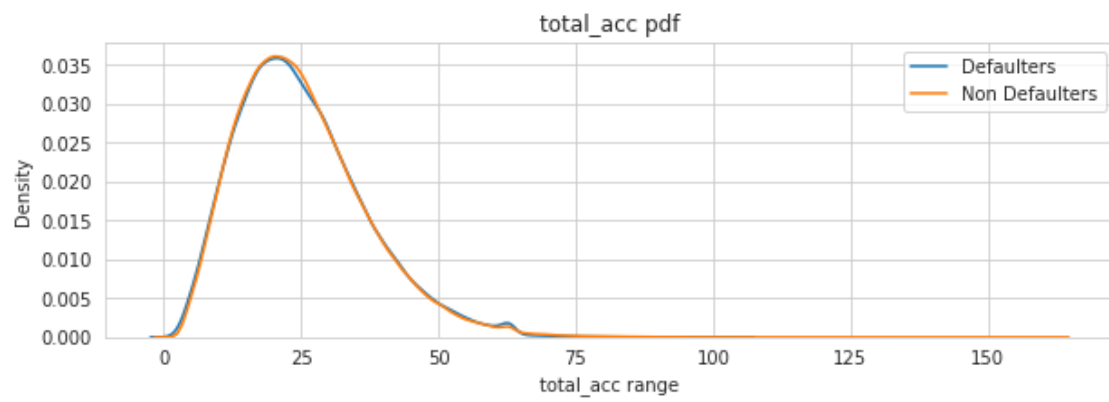
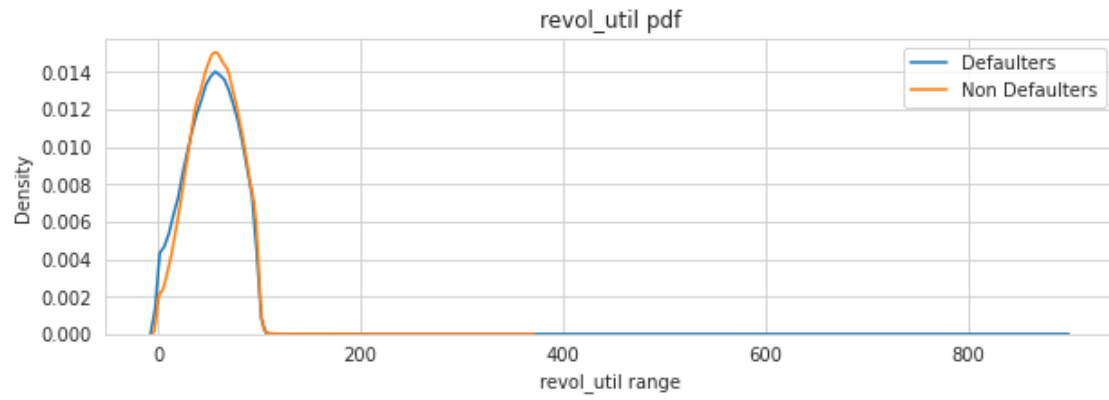
```
[ ]: for idx,cat_col in enumerate(preprocess_int_feature):
      plotPdf(cat_col)
```

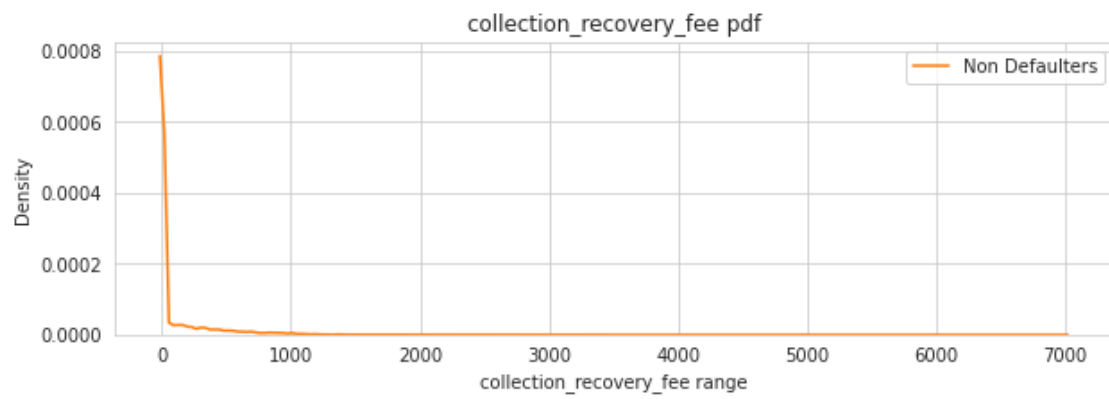
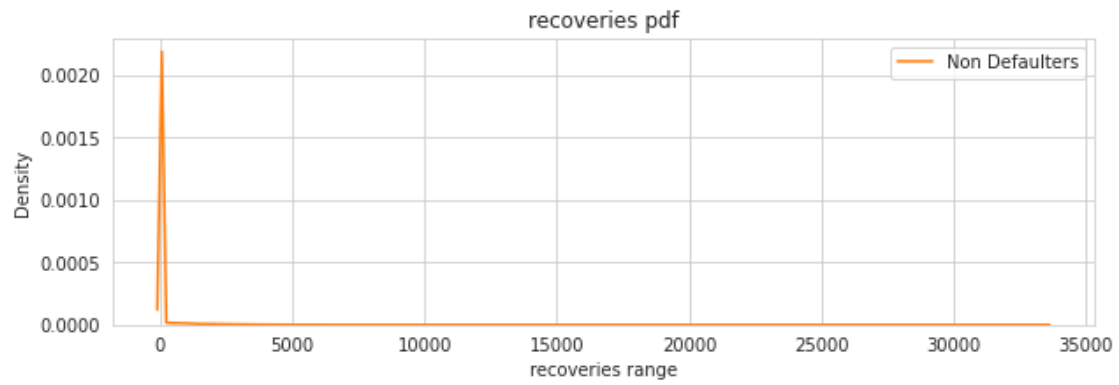
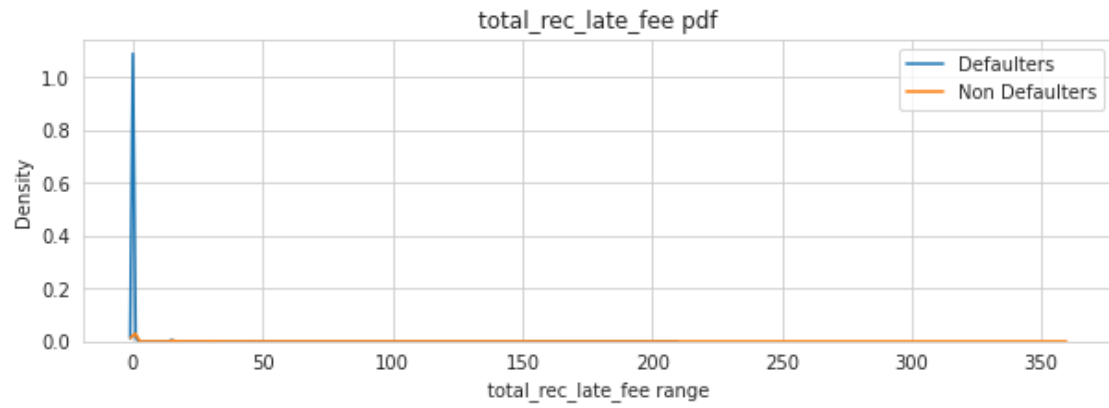


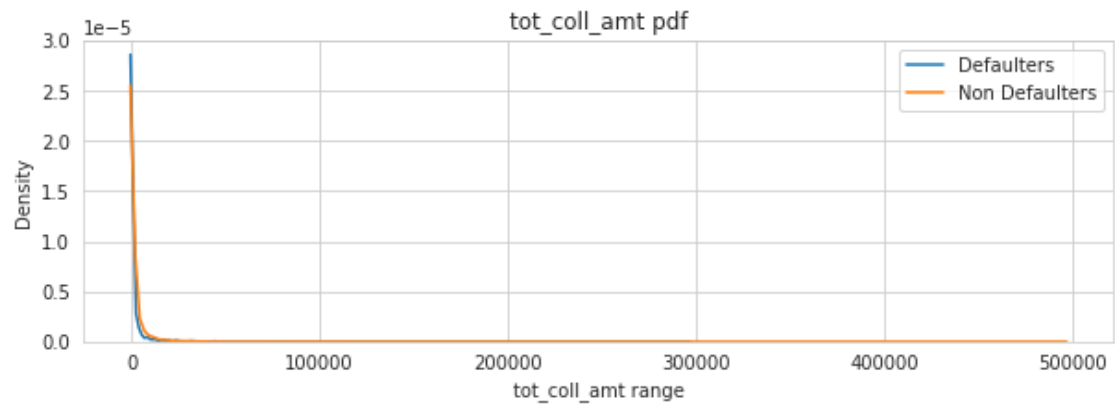
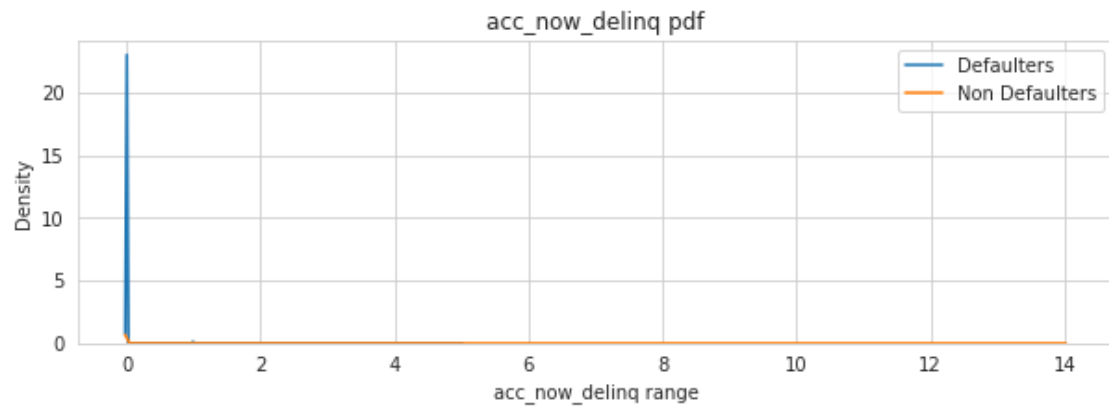
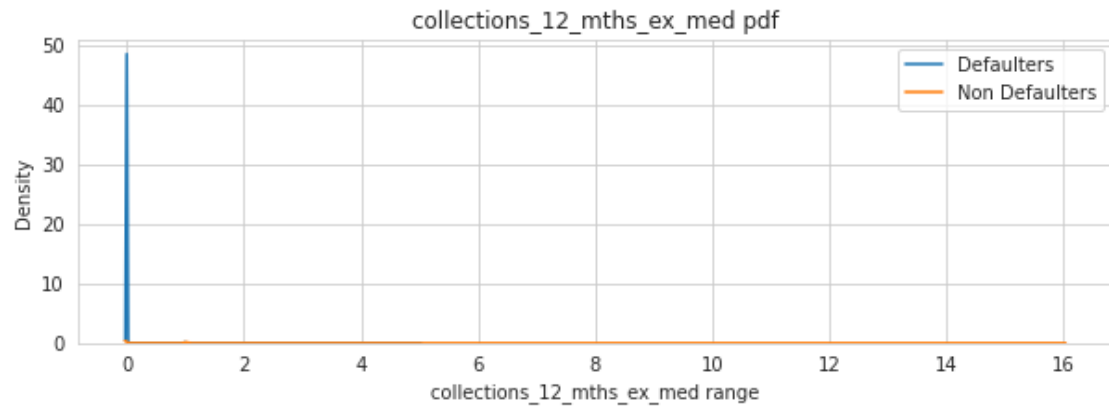



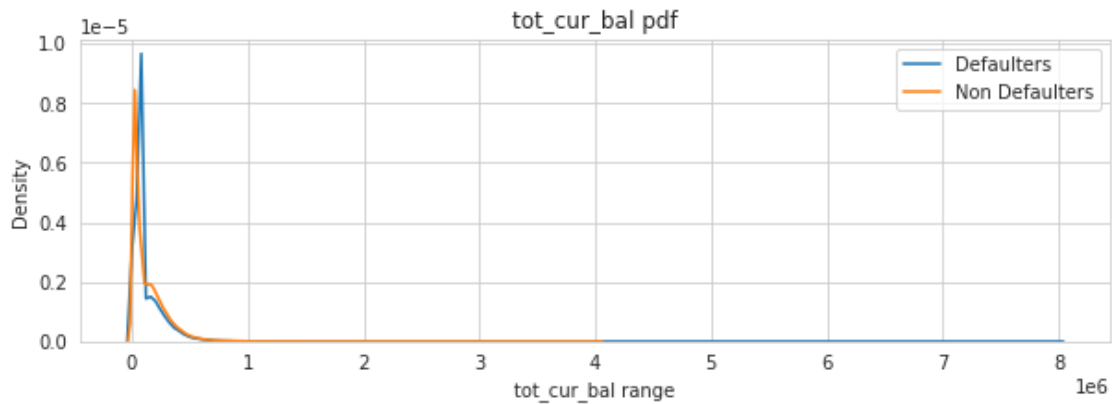












```
[ ]: train_data_preprocessed[train_data_preprocessed.loan_status == 1].
      ↪collection_recovery_fee.value_counts()
```

```
[ ]: 0.0    125827
      Name: collection_recovery_fee, dtype: int64
```

```
[ ]: train_data_preprocessed[train_data_preprocessed.loan_status == 1].recoveries.
      ↪value_counts()
```

```
[ ]: 0.0    125827
      Name: recoveries, dtype: int64
```

- from above analysis we can say that 'loan_amnt', 'funded_amnt', 'funded_amnt_inv', 'int_rate' these are some imp features because they show quite good separation of defaulter and no defaulter pdf areas
- collection_recovery_fee and recoveries having zero std deviation for defaulters i.e all values for defaulter cases for this two features are zero

3 create new feature

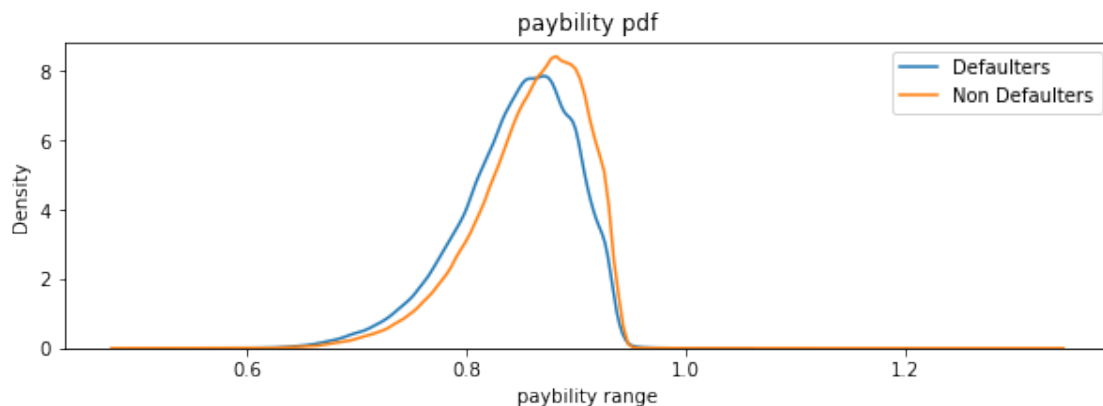
- we dont have experties in bank domain , generally we can see if there is to much of loan with respect to monthly income then person will no able to complete loan
- new feature = $\log(\text{loan_amount} + \text{intrest_rate} * \text{term_of_loan}) + 0.00001 / \log(\text{anual_income}) + 0.00001$
- used log to reduce correlaiton and 0.00001 noice prevent us from devide by zero error
- if pdf of this feature will not show significant seperation then we will take this feature

```
[ ]: import math
train_data_preprocessed['paybility'] = train_data_preprocessed.apply(lambda row:
    ↳ (math.log(row.loan_amnt + (float(row.term.split('_')[0]) * row.int_rate)) +
    ↳ 0.00001 )/math.log(row.annual_inc), axis=1)
```

```
[ ]: train_data_preprocessed['paybility']
```

```
[ ]: 0          0.937054
      1          0.772015
      2          0.862030
      3          0.835694
      4          0.893638
      ...
532423    0.884233
532424    0.861602
532425    0.929137
532426    0.924028
532427    0.848782
Name: paybility, Length: 532428, dtype: float64
```

```
[ ]: plotPdf('paybility')
```



- its seperating regions cosiderable we will keep this feature

4 Modeling

- train test split
- oversampling of train dataset
- normalize numerical features as we dont want to loose original distribution of dataser
- one hot encoading of categorical features
- after dataset is ready load into model we will apply (Logistic regression , GBDT , Random-Forest)
 - cross validation using random serach
 - plot auc-roc value against epocs
 - draw roc cusrve on both train and test
 - draw confussion metrix in from best threshold
- we will also try small mlp nural network with dence layers
- compaire models and select model having heighest AUC-ROC score
- predict submission data

```
[ ]: #train_data_preprocessed.to_pickle("train_data_preprocessed.pkl")
train_data_preprocessed = pd.read_pickle("train_data_preprocessed.pkl")
```

```
[ ]: train_data_preprocessed.shape
```

```
[ ]: (532428, 36)
```

```
[ ]: y = train_data_preprocessed['loan_status']
X_data = train_data_preprocessed.drop(columns=['loan_status'])
```

```
[ ]: from sklearn.model_selection import train_test_split
X_train_old, X_test, y_train_old, y_test = train_test_split(X_data, y,
                                                            stratify=y,
                                                            test_size=0.2,
                                                            random_state=42)
```

```
[ ]: #!pip install -U imbalanced-learn
from imblearn.over_sampling import RandomOverSampler
from collections import Counter
```

```
[ ]: print(Counter(y_train_old))
# define oversampling strategy
oversample = RandomOverSampler(sampling_strategy=0.5, random_state=42)
# fit and apply the transform
X_over, y_over = oversample.fit_resample(X_train_old, y_train_old)
# summarize class distribution
print(Counter(y_over))
```

```
Counter({0: 325280, 1: 100662})
```

```
Counter({0: 325280, 1: 162640})
```



```
[ ]: X_train , y_train = X_over, y_over
```

5 normalizing numerical features

```
[ ]: from sklearn.preprocessing import Normalizer

numerical_features = ['loan_amnt', 'int_rate',
                      'annual_inc', 'dti', 'delinq_2yrs', 'inq_last_6mths', 'open_acc',
                      'pub_rec', 'revol_bal', 'revol_util', 'total_acc', 'total_rec_int',
                      'total_rec_late_fee', 'recoveries', 'collection_recovery_fee',
                      'collections_12_mths_ex_med', 'acc_now_delinq', 'tot_coll_amt',
                      'tot_cur_bal', 'total_rev_hi_lim', 'payability']
```

```
[ ]: X_test.shape
```

```
[ ]: (106486, 35)
```

```
[ ]: X_train.shape
```

```
[ ]: (487920, 35)
```

```
[ ]: X_train_numerical_features = X_train[numerical_features].values
     X_test_numerical_features = X_test[numerical_features].values
```

```
[ ]: X_train[numerical_features]
```

```
[ ]:
```

	loan_amnt	int_rate	...	total_rev_hi_lim	payability
0	8950	15.99	...	21100.0	0.904719
1	11900	18.54	...	15900.0	0.919721
2	6000	12.49	...	15700.0	0.810718
3	5200	6.54	...	23700.0	0.766193
4	10675	20.99	...	5000.0	0.911511
...
487915	10000	5.99	...	23700.0	0.861016
487916	8000	10.99	...	50100.0	0.791277
487917	21550	15.61	...	49700.0	0.867431
487918	15000	12.99	...	21200.0	0.854447
487919	10000	11.49	...	23700.0	0.865844

```
[487920 rows x 21 columns]
```

```
[ ]: from sklearn.preprocessing import MinMaxScaler
     def normalize_feature(nnumerical_fetures):
         X_train_numerical_features = X_train[nnumerical_fetures].values
         X_test_numerical_features = X_test[nnumerical_fetures].values
```

```

data = X_train_numerical_features
scaler = MinMaxScaler()
scaler.fit(data)
pickle.dump(scaler, open('normalizers_numeric.pkl', 'wb'))
X_train_numerical = scaler.transform(X_train_numerical_features)
X_test_numerical = scaler.transform(X_test_numerical_features)
return X_train_numerical , X_test_numerical

```

```
[ ]: X_train_numerical , X_test_numerical = normalize_feature(numerical_features)
```

```
[ ]: X_train_numerical.shape
```

```
[ ]: (487920, 21)
```

```
[ ]: X_train.loan_amnt.value_counts()
```

```

[ ]: 10000    34360
      12000    27810
      15000    25783
      20000    25199
      35000    19220
      ...
      34700         1
      29075         1
       925          1
      32575         1
      33725         1
      Name: loan_amnt, Length: 1368, dtype: int64

```

```

[ ]: from sklearn.preprocessing import OneHotEncoder
def vectorize_feature(features):
    for i in features:
        vectorizer = OneHotEncoder(handle_unknown='ignore')
        vectorizer.fit(X_train[i].values.reshape(-1,1)) # fit has to happen only on
        ↪ train data
        globals()[f"vectorizer_{i}"] = vectorizer
        pickle.dump(globals()[f"vectorizer_{i}"], open(f'vectorizers/vectorizer_{i}.
        ↪ pkl', 'wb'))
        # we use the fitted OneHotEncoder to convert the text to vector
        X_train_vectors = vectorizer.transform(X_train[i].values.reshape(-1,1))
        X_test_vectors = vectorizer.transform(X_test[i].values.reshape(-1,1))

        globals()[f"X_train_{i}"] = X_train_vectors
        globals()[f"X_test_{i}"] = X_test_vectors
        print("After one hot encoding of "+i)
        print(X_train_vectors.shape, y_train.shape)
        print(X_test_vectors.shape, y_test.shape)

```

```
print("="*100)
```

```
[ ]: categorical_features_preprocessed =  
    ↳ ['term', 'batch_enrolled', 'grade', 'sub_grade', 'emp_title', 'emp_length'  
    , 'home_ownership', 'verification_status', 'purpose', 'title'  
    , 'zip_code', 'addr_state', 'initial_list_status'  
    , 'last_week_pay']  
len(categorical_features_preprocessed)
```

```
[ ]: 14
```

```
[ ]: vectorize_feature(categorical_features_preprocessed)
```

After one hot encoding of term

(487920, 2) (487920,)

(106486, 2) (106486,)

=====

=====

After one hot encoding of batch_enrolled

(487920, 105) (487920,)

(106486, 105) (106486,)

=====

=====

After one hot encoding of grade

(487920, 7) (487920,)

(106486, 7) (106486,)

=====

=====

After one hot encoding of sub_grade

(487920, 35) (487920,)

(106486, 35) (106486,)

=====

=====

After one hot encoding of emp_title

(487920, 138934) (487920,)

(106486, 138934) (106486,)

=====

=====

After one hot encoding of emp_length

(487920, 12) (487920,)

(106486, 12) (106486,)

=====

=====

After one hot encoding of home_ownership

(487920, 6) (487920,)

(106486, 6) (106486,)

=====

```

=====
After one hot encoding of verification_status
(487920, 3) (487920,)
(106486, 3) (106486,)
=====
=====
After one hot encoding of purpose
(487920, 14) (487920,)
(106486, 14) (106486,)
=====
=====
After one hot encoding of title
(487920, 27853) (487920,)
(106486, 27853) (106486,)
=====
=====
After one hot encoding of zip_code
(487920, 907) (487920,)
(106486, 907) (106486,)
=====
=====
After one hot encoding of addr_state
(487920, 51) (487920,)
(106486, 51) (106486,)
=====
=====
After one hot encoding of initial_list_status
(487920, 2) (487920,)
(106486, 2) (106486,)
=====
=====
After one hot encoding of last_week_pay
(487920, 97) (487920,)
(106486, 97) (106486,)
=====
=====

```

```

[ ]: from scipy.sparse import hstack
X_train_final = hstack((X_train_numerical , X_train_term_
    ↪,X_train_batch_enrolled,X_train_grade,X_train_sub_grade,X_train_emp_title,X_train_emp_lengt
,X_train_home_ownership,X_train_verification_status,X_train_purpose,X_train_title
,X_train_zip_code,X_train_addr_state,X_train_initial_list_status
,X_train_last_week_pay ))

[ ]: pickle.dump(X_train_final, open('X_train_final.pkl', 'wb'))

```

```
[ ]: X_test_final = hstack((X_test_numerical , X_test_term
    ↪,X_test_batch_enrolled,X_test_grade,X_test_sub_grade,X_test_emp_title,X_test_emp_length
    ,X_test_home_ownership,X_test_verification_status,X_test_purpose,X_test_title
    ,X_test_zip_code,X_test_addr_state,X_test_initial_list_status
    ,X_test_last_week_pay ))
pickle.dump(X_test_final, open('X_test_final.pkl', 'wb'))

[ ]: pickle.dump(y_train, open('y_train.pkl', 'wb'))
pickle.dump(y_test, open('y_test.pkl', 'wb'))

[ ]: pickle.dump(X_train_numerical, open('X_train_numerical.pkl', 'wb'))
pickle.dump(X_test_numerical, open('X_test_numerical.pkl', 'wb'))

[ ]: X_train_final = pickle.load(open('X_train_final.pkl', 'rb'))
X_test_final = pickle.load(open('X_test_final.pkl', 'rb'))

[ ]: X_train_csr = X_train_final.tocsr()
X_test_csr = X_test_final.tocsr()

[ ]: y_train = pickle.load(open('y_train.pkl', 'rb'))
y_test = pickle.load(open('y_test.pkl', 'rb'))

[ ]: X_test_final

[ ]: <106486x168049 sparse matrix of type '<class 'numpy.float64'>'
    with 2836048 stored elements in COOrdinate format>
```

#Apply Model

- first we apply logistic regression because Liniar model give good results when dimentions are heigh

```
[ ]: def batch_predict(clf, data):
    """ this function predicts batch wise propablity
        for given data over given model
        """

    y_data_pred = []
    tr_loop = data.shape[0] - data.shape[0]%1000
    for i in range(0, tr_loop, 1000):
        y_data_pred.extend(clf.predict_proba(data[i:i+1000]))[:,1])
    # we will be predicting for the last data points
    if data.shape[0]%1000 !=0:
        y_data_pred.extend(clf.predict_proba(data[tr_loop:]))[:,1])

    return y_data_pred
```

```
[ ]: from sklearn.model_selection import RandomizedSearchCV
```

Logistic Regression

```
[ ]: from sklearn.datasets import load_iris
from sklearn.linear_model import LogisticRegression
clf = LogisticRegression(random_state=15)
parameters = {'C':[1e-7,1e-6,1e-5, 1e-4, 1e-3, 1e-2, 1e-1, 1, 10, 100]}
clf = RandomizedSearchCV(clf, parameters, cv=3 ,scoring='roc_auc',
    ↪,return_train_score=True,random_state=15)

clf.fit(X_train_csr,y_train)

results = pd.DataFrame.from_dict(clf.cv_results_)
print(results.columns)
results = results.sort_values(['param_C'])

train_auc= results['mean_train_score']
train_auc_std= results['std_train_score']
cv_auc = results['mean_test_score']
cv_auc_std= results['std_test_score']
K = list(map(lambda x: np.log(x) , results['param_C']))
#K = results['param_C']
plt.plot(K, train_auc, label='Train AUC')

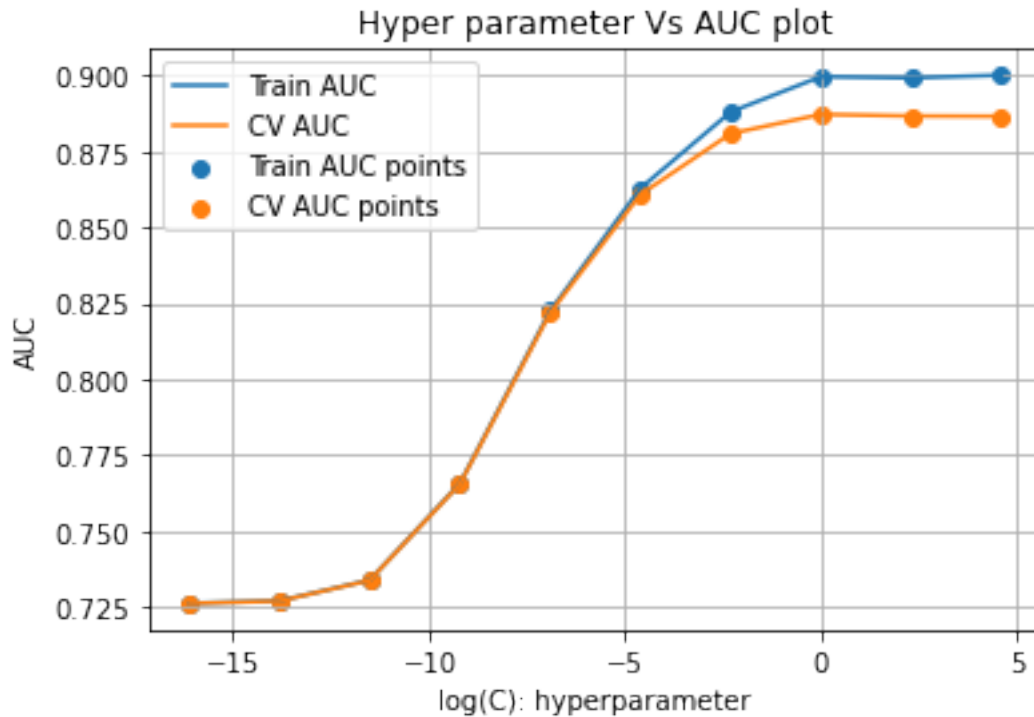
plt.plot(K, cv_auc, label='CV AUC')

plt.scatter(K, train_auc, label='Train AUC points')
plt.scatter(K, cv_auc, label='CV AUC points')

plt.legend()
plt.xlabel("log(C): hyperparameter")
plt.ylabel("AUC")
plt.title("Hyper parameter Vs AUC plot")
plt.grid()
plt.show()

results
```

```
Index(['mean_fit_time', 'std_fit_time', 'mean_score_time', 'std_score_time',
      'param_C', 'params', 'split0_test_score', 'split1_test_score',
      'split2_test_score', 'mean_test_score', 'std_test_score',
      'rank_test_score', 'split0_train_score', 'split1_train_score',
      'split2_train_score', 'mean_train_score', 'std_train_score'],
      dtype='object')
```



```
[ ]:  mean_fit_time  std_fit_time  ...  mean_train_score  std_train_score
0      1.931660    0.041759  ...      0.726235      0.000389
1      1.559254    0.073042  ...      0.727051      0.000382
2      1.967143    0.039396  ...      0.733847      0.000328
3      3.106050    0.016808  ...      0.765705      0.000086
4      6.956268    0.078847  ...      0.822467      0.000044
5     15.453114    0.261965  ...      0.862777      0.000250
6     15.177733    0.358979  ...      0.887906      0.001693
7     15.282248    0.173678  ...      0.899523      0.001926
8     15.434415    0.144933  ...      0.899126      0.001063
9     15.478861    0.116143  ...      0.899997      0.002875
```

[10 rows x 17 columns]

```
[ ]:
```

```
[ ]: #best C is 1
from sklearn.metrics import roc_curve, auc

lr_model = LogisticRegression(C=1, random_state=15)
lr_model.fit(X_train_csr, y_train)
# roc_auc_score(y_true, y_score) the 2nd parameter should be probability
↳ estimates of the positive class
```

```

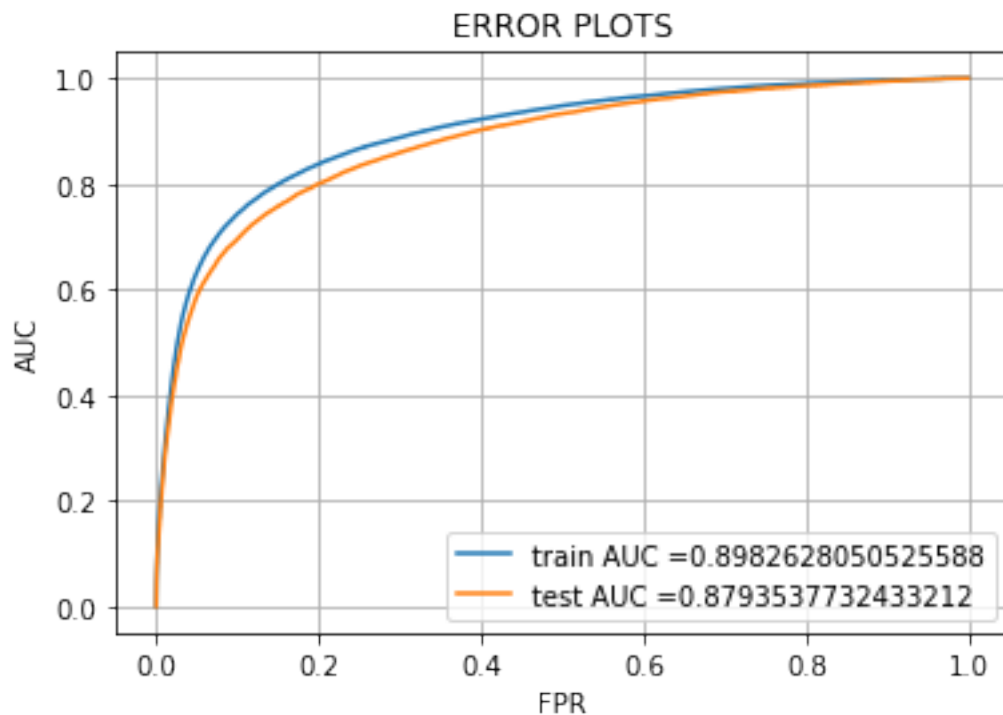
# not the predicted outputs

y_train_pred = batch_predict(lr_model, X_train_csr)
y_test_pred = batch_predict(lr_model, X_test_csr)

train_fpr, train_tpr, tr_thresholds = roc_curve(y_train, y_train_pred)
test_fpr, test_tpr, te_thresholds = roc_curve(y_test, y_test_pred)

plt.plot(train_fpr, train_tpr, label="train AUC =" + str(auc(train_fpr,
    ↪ train_tpr)))
plt.plot(test_fpr, test_tpr, label="test AUC =" + str(auc(test_fpr, test_tpr)))
plt.legend()
plt.xlabel("FPR")
plt.ylabel("AUC")
plt.title("ERROR PLOTS")
plt.grid()
plt.show()
pickle.dump(lr_model, open('lr_model', 'wb'))

```



```

[ ]: def find_best_threshold(threshold, fpr, tpr):
    t = threshold[np.argmax(tpr*(1-fpr))]
    # (tpr*(1-fpr)) will be maximum if your fpr is very low and tpr is very high

```



```

    print("the maximum value of tpr*(1-fpr)", max(tpr*(1-fpr)), "for_
↪threshold", np.round(t,3))
    return t

def predict_with_best_t(proba, threshold):
    predictions = []
    for i in proba:
        if i>=threshold:
            predictions.append(1)
        else:
            predictions.append(0)
    return predictions

```

```

[ ]: import seaborn as sns
import matplotlib.pyplot as plt
print("="*100)
from sklearn.metrics import confusion_matrix
best_t = find_best_threshold(tr_thresholds, train_fpr, train_tpr)
print("Train confusion matrix")
train_predict = predict_with_best_t(y_train_pred, best_t)
cm_train = confusion_matrix(y_train, predict_with_best_t(y_train_pred, best_t))
#cm_train = confusion_matrix(y_train, train_predict)
#print(confusion_matrix(y_train, predict_with_best_t(y_train_pred, best_t)))
print("Test confusion matrix")
test_predict = predict_with_best_t(y_test_pred, best_t)
cm_test = confusion_matrix(y_test, predict_with_best_t(y_test_pred, best_t))
#cm_test = confusion_matrix(y_test, test_predict)
#print(confusion_matrix(y_test, predict_with_best_t(y_test_pred, best_t)))

fig = plt.figure(figsize = (5,5)) # width x height
ax1 = fig.add_subplot(2,2,1) # row, column, position
ax2 = fig.add_subplot(2,2,3)


ax1.set_title('training set')
cm_train = pd.DataFrame(cm_train)
cm_train['Predicted'] = [0,1]
cm_train.columns.name = 'Actual'
cm_train.set_index('Predicted',inplace=True)
sns.heatmap(cm_train,ax=ax1,square=True, annot=True,fmt="d",cmap='Blues')

ax2.set_title('testing set')
cm_test = pd.DataFrame(cm_test)
cm_test['Predicted'] = [0,1]
cm_test.columns.name = 'Actual'
cm_test.set_index('Predicted',inplace=True)
sns.heatmap(cm_test,ax=ax2,square=True, annot=True,fmt="d",cmap='Blues')

```

```
=====
the maximum value of tpr*(1-fpr) 0.679641655354506 for threshold 0.335
Train confusion matrix
Test confusion matrix
```

```
[ ]: <matplotlib.axes._subplots.AxesSubplot at 0x7f277ff33250>
```



5.1 GBDT

```
[ ]: import xgboost as xgb
from xgboost.sklearn import XGBClassifier
```

```
[ ]: parameters = {
    'learning_rate': [0.05, 0.1, 0.2, 0.3],
    'n_estimators': [10, 50, 120, 200]
}
xgb = XGBClassifier(
    max_depth=3,
    subsample=0.8,
```

```

colsample_bytree=0.8)

random_search =
    ↳RandomizedSearchCV(xgb,parameters,cv=3,return_train_score=True,random_state=15,scoring='roc')
random_search.fit(X_train_csr,y_train)

results = pd.DataFrame.from_dict(random_search.cv_results_)

results

```

```

[ ]:   mean_fit_time  std_fit_time  ...  mean_train_score  std_train_score
0      9.667801      0.496915  ...      0.764585      0.001853
1     31.940976      0.431594  ...      0.860757      0.004607
2    117.651682      1.100278  ...      0.858403      0.000115
3     71.379337      0.822498  ...      0.838266      0.000358
4     71.361843      0.580145  ...      0.907090      0.003358
5     32.100068      0.288634  ...      0.798859      0.001024
6    117.399010      1.028311  ...      0.938808      0.001881
7     71.746024      0.644654  ...      0.865048      0.000668
8     31.915615      0.393351  ...      0.879740      0.003140
9    116.421143      1.177270  ...      0.926831      0.002243

```

[10 rows x 18 columns]

```

[ ]: import plotly.offline as offline
import plotly.graph_objs as go
offline.init_notebook_mode()
import numpy as np
import matplotlib.pyplot as plt

```

5.1.1 please ensure you have installed plotly to see below 3d graph

```

[ ]: x1,y1,z1 = results['param_n_estimators'],results['param_learning_rate'],
    ↳results['mean_train_score']
x2,y2,z2 = results['param_n_estimators'],results['param_learning_rate'],
    ↳results['mean_test_score']
trace1 = go.Scatter3d(x=x1,y=y1,z=z1, name = 'train')
trace2 = go.Scatter3d(x=x2,y=y2,z=z2, name = 'Cross validation')
data = [trace1, trace2]

layout = go.Layout(scene = dict(
    xaxis = dict(title='param_n_estimators'),
    yaxis = dict(title='param_learning_rate'),
    zaxis = dict(title='AUC'),))

```

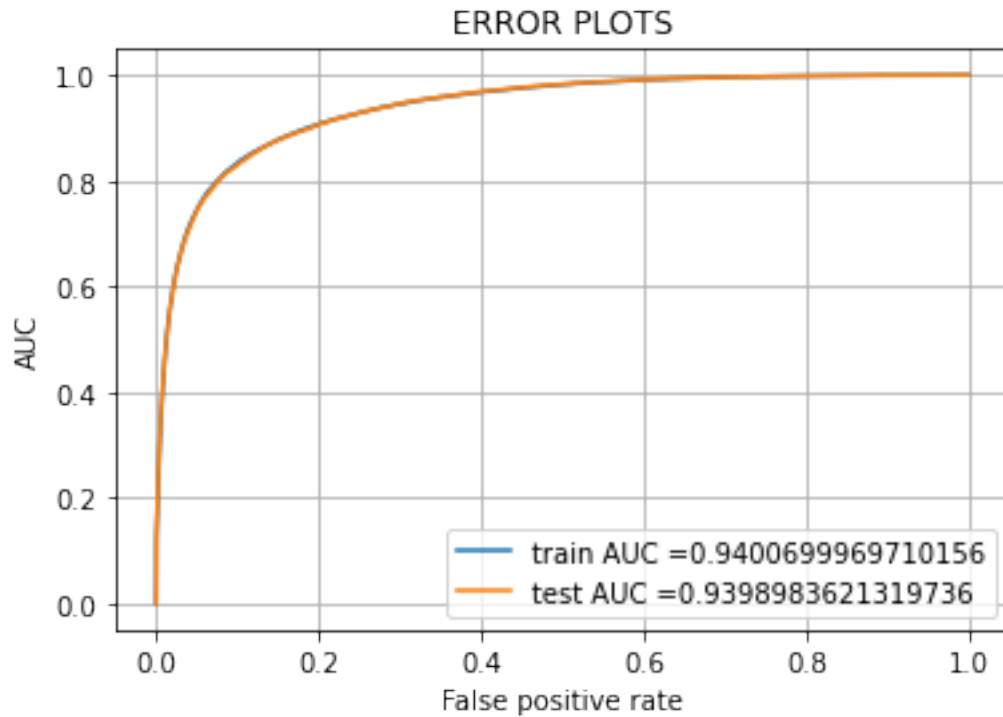
```
fig = go.Figure(data=data, layout=layout)
offline.iplot(fig, filename='3d-scatter-colorscale')
fig.show(renderer="colab")
plt.show()
```

```
[ ]: # {'n_estimators': 200, 'learning_rate': 0.3} these are our best params
from sklearn.metrics import roc_curve, auc
xgbClf = XGBClassifier(
    learning_rate = 0.3,
    n_estimators=200,
    max_depth=3,
    subsample=0.8,
    colsample_bytree=0.8)
xgbClf.fit(X_train_csr, y_train)
# roc_auc_score(y_true, y_score) the 2nd parameter should be probability
# → estimates of the positive class
# not the predicted outputs

y_train_pred = batch_predict(xgbClf, X_train_csr)
y_test_pred = batch_predict(xgbClf, X_test_csr)

train_fpr, train_tpr, tr_thresholds = roc_curve(y_train, y_train_pred)
test_fpr, test_tpr, te_thresholds = roc_curve(y_test, y_test_pred)

plt.plot(train_fpr, train_tpr, label="train AUC =" + str(auc(train_fpr,
    → train_tpr)))
plt.plot(test_fpr, test_tpr, label="test AUC =" + str(auc(test_fpr, test_tpr)))
plt.legend()
plt.xlabel("False positive rate")
plt.ylabel("AUC")
plt.title("ERROR PLOTS")
plt.grid()
plt.show()
```



```
[ ]: from sklearn.externals import joblib
```

```
joblib_file = "xgb_model.joblib"
joblib.dump(xgbClf, joblib_file)
```

```
[ ]: ['xgb_model.joblib']
```

```
[ ]: xgbClf = joblib.load('xgb_model.joblib')
```

```
[ ]: print("="*100)
best_t = find_best_threshold(tr_thresholds, train_fpr, train_tpr)
print("Train confusion matrix")
train_predict = predict_with_best_t(y_train_pred, best_t)
cm_train = confusion_matrix(y_train, predict_with_best_t(y_train_pred, best_t))
#cm_train = confusion_matrix(y_train, train_predict)
#print(confusion_matrix(y_train, predict_with_best_t(y_train_pred, best_t)))
print("Test confusion matrix")
test_predict = predict_with_best_t(y_test_pred, best_t)
cm_test = confusion_matrix(y_test, predict_with_best_t(y_test_pred, best_t))
#cm_test = confusion_matrix(y_test, test_predict)
#print(confusion_matrix(y_test, predict_with_best_t(y_test_pred, best_t)))

fig = plt.figure(figsize = (5,5)) # width x height
```

```

ax1 = fig.add_subplot(2,2,1) # row, column, position
ax2 = fig.add_subplot(2,2,3)

ax1.set_title('training set')
cm_train = pd.DataFrame(cm_train)
cm_train['Predicted'] = [0,1]
cm_train.columns.name = 'Actual'
cm_train.set_index('Predicted',inplace=True)
sns.heatmap(cm_train,ax=ax1,square=True, annot=True,fmt="d",cmap='Blues')

ax2.set_title('testing set')
cm_test = pd.DataFrame(cm_test)
cm_test['Predicted'] = [0,1]
cm_test.columns.name = 'Actual'
cm_test.set_index('Predicted',inplace=True)
sns.heatmap(cm_test,ax=ax2,square=True, annot=True,fmt="d",cmap='Blues',)

```


=====

the maximum value of $tpr \cdot (1 - fpr)$ 0.7521867978163548 for threshold 0.342

Train confusion matrix

Test confusion matrix

[]: <matplotlib.axes._subplots.AxesSubplot at 0x7f5c2b0d7990>



5.2 Random Forest

```
[ ]: from sklearn.ensemble import RandomForestClassifier
parameters={'n_estimators':[10,30,50,80,100],'max_depth': [1, 5, 10, 50]}
clf_tree=RandomForestClassifier()
random_search =_
↳RandomizedSearchCV(clf_tree,parameters,cv=3,return_train_score=True,random_state=15,n_iter=
random_search.fit(X_train_csr,y_train)

results = pd.DataFrame.from_dict(random_search.cv_results_)
print(results.columns)
results = results.sort_values(['param_max_depth'])

results
```

```
Index(['mean_fit_time', 'std_fit_time', 'mean_score_time', 'std_score_time',
      'param_n_estimators', 'param_max_depth', 'params', 'split0_test_score',
      'split1_test_score', 'split2_test_score', 'mean_test_score',
      'std_test_score', 'rank_test_score', 'split0_train_score',
      'split1_train_score', 'split2_train_score', 'mean_train_score',
      'std_train_score'],
      dtype='object')
```

```
[ ]: mean_fit_time  std_fit_time  ...  mean_train_score  std_train_score
2          8.053658    0.755660  ...          0.710498         0.009089
4          9.820057    0.928637  ...          0.691998         0.032478
5          5.109497    0.477210  ...          0.637429         0.010957
9          3.221154    0.271778  ...          0.618996         0.018599
3         35.680916    4.428566  ...          0.791857         0.008366
10         10.799898    1.211847  ...          0.721833         0.005188
13         17.858352    2.268550  ...          0.757287         0.008296
6          7.153913    0.885289  ...          0.724137         0.013648
8         54.829779    6.657790  ...          0.809414         0.006388
12         68.294789    8.097947  ...          0.808040         0.006578
14         21.078361    2.516116  ...          0.788913         0.005481
0         452.040343   38.835618  ...          0.900459         0.004362
1          44.561542    4.798198  ...          0.848499         0.010950
7         134.166251   14.245967  ...          0.883674         0.003985
11         224.039160   21.466611  ...          0.892269         0.004467
15         363.562780   37.001578  ...          0.897993         0.003660
```

```
[16 rows x 18 columns]
```

```
[ ]:
```

```
[ ]: x1,y1,z1 = results['param_n_estimators'],results['param_max_depth'],
      ↳results['mean_train_score']
x2,y2,z2 = results['param_n_estimators'],results['param_max_depth'],
      ↳results['mean_test_score']
trace1 = go.Scatter3d(x=x1,y=y1,z=z1, name = 'train')
trace2 = go.Scatter3d(x=x2,y=y2,z=z2, name = 'Cross validation')
data = [trace1, trace2]

layout = go.Layout(scene = dict(
    xaxis = dict(title='param_n_estimators'),
    yaxis = dict(title='param_max_depth'),
    zaxis = dict(title='AUC'),))

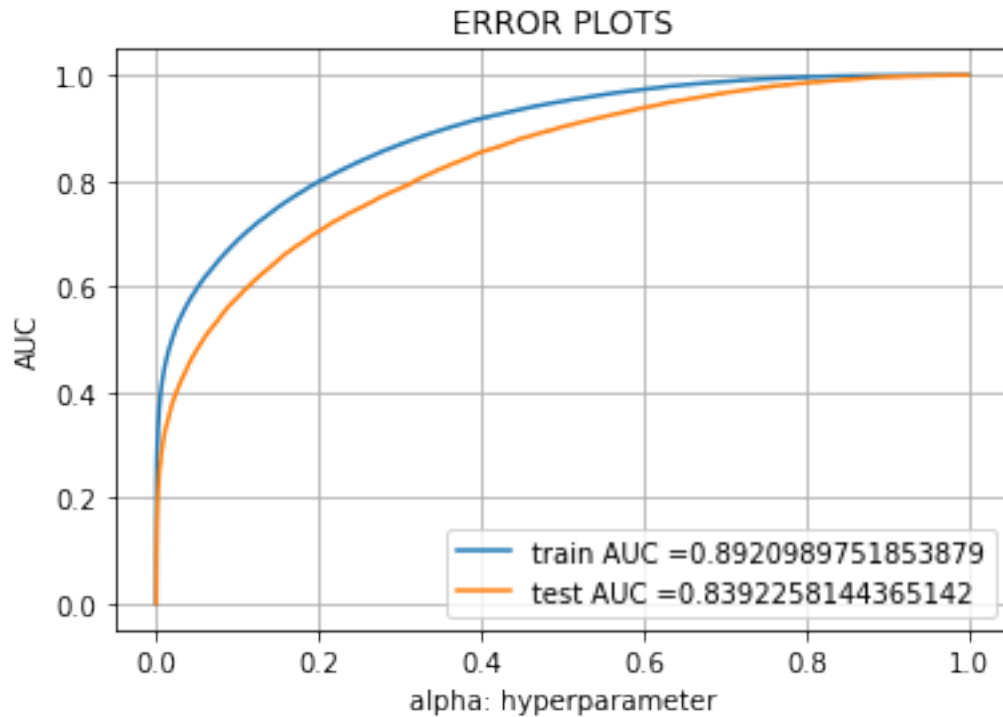
fig = go.Figure(data=data, layout=layout)
offline.iplot(fig, filename='3d-scatter-colorscale')
fig.show(renderer="colab")
plt.show()

[ ]: #{'n_estimators': 100, 'max_depth': 50} are the best parameters
dtree = RandomForestClassifier(n_estimators=100,max_depth=50,random_state=15)
dtree.fit(X_train_csr, y_train)
# roc_auc_score(y_true, y_score) the 2nd parameter should be probability
↳estimates of the positive class
# not the predicted outputs

y_train_pred = batch_predict(dtree, X_train_csr)
y_test_pred = batch_predict(dtree, X_test_csr)

train_fpr, train_tpr, tr_thresholds = roc_curve(y_train, y_train_pred)
test_fpr, test_tpr, te_thresholds = roc_curve(y_test, y_test_pred)

plt.plot(train_fpr, train_tpr, label="train AUC =" +str(auc(train_fpr,
↳train_tpr)))
plt.plot(test_fpr, test_tpr, label="test AUC =" +str(auc(test_fpr, test_tpr)))
plt.legend()
plt.xlabel("alpha: hyperparameter")
plt.ylabel("AUC")
plt.title("ERROR PLOTS")
plt.grid()
plt.show()
```

```
[ ]: print("="*100)
from sklearn.metrics import confusion_matrix
best_t = find_best_threshold(tr_thresholds, train_fpr, train_tpr)
print("Train confusion matrix w2v")
train_predict = predict_with_best_t(y_train_pred, best_t)
cm_train = confusion_matrix(y_train, predict_with_best_t(y_train_pred, best_t))
#cm_train = confusion_matrix(y_train, train_predict)
#print(confusion_matrix(y_train, predict_with_best_t(y_train_pred, best_t)))
print("Test confusion matrix w2v")
test_predict = predict_with_best_t(y_test_pred, best_t)
cm_test = confusion_matrix(y_test, predict_with_best_t(y_test_pred, best_t))
#cm_test = confusion_matrix(y_test, test_predict)
#print(confusion_matrix(y_test, predict_with_best_t(y_test_pred, best_t)))

fig = plt.figure(figsize = (5,5)) # width x height
ax1 = fig.add_subplot(2,2,1) # row, column, position
ax2 = fig.add_subplot(2,2,3)

ax1.set_title('training set')
cm_train = pd.DataFrame(cm_train)
cm_train['Predicted'] = [0,1]
cm_train.columns.name = 'Actual'
cm_train.set_index('Predicted',inplace=True)
```

```

sns.heatmap(cm_train,ax=ax1,square=True, annot=True,fmt="d",cmap='Blues')

ax2.set_title('testing set')
cm_test = pd.DataFrame(cm_test)
cm_test['Predicted'] = [0,1]
cm_test.columns.name = 'Actual'
cm_test.set_index('Predicted',inplace=True)
sns.heatmap(cm_test,ax=ax2,square=True, annot=True,fmt="d",cmap='Blues')

```


```

=====
=====

```

the maximum value of $tpr \cdot (1 - fpr)$ 0.6399008804688818 for threshold 0.341
Train confusion matrix w2v
Test confusion matrix w2v

```
[ ]: <matplotlib.axes._subplots.AxesSubplot at 0x7f5c2c488950>
```



```
[ ]: joblib.dump(dtree, 'rdtree.joblib')
```

```
[ ]: ['rdtree.joblib']
```

```
[ ]: import tensorflow as tf
import random as rn
from sklearn.model_selection import train_test_split
from tensorflow.keras.callbacks import Callback
from sklearn.metrics import confusion_matrix, f1_score, precision_score,
    ↪recall_score
from tensorflow.keras.callbacks import ModelCheckpoint
from tensorflow.keras.callbacks import EarlyStopping
import tensorflow.keras.layers as layers
from tensorflow.keras.layers import BatchNormalization
from tensorflow.keras.layers import Dense, Input, Dropout
from tensorflow.keras.models import Model
```

```
[ ]: def aucroc(y_true, y_pred):
    m = tf.keras.metrics.AUC()
    m.update_state(y_true, y_pred)
    return m.result().numpy()
```

```
[ ]: def chage_sparse_to_dence(sparce_mat):
    reaturn_mat = []
    for i in tqdm(range(0, sparce_mat.shape[0])):
        if reaturn_mat == []:
            reaturn_mat = sparce_mat[i,:].toarray()
            continue
        reaturn_mat = np.concatenate((reaturn_mat, sparce_mat[i,:].toarray() ),
    ↪axis=0)
    return reaturn_mat
```

```
[ ]: X_train_dence = pd.DataFrame.sparse.from_spmatrix(X_train_csr)
X_test_dence = pd.DataFrame.sparse.from_spmatrix(X_test_csr)
```

```
[ ]: X_train_dence
```

```
[ ]:
      0      1      2      3      ...  168045  168046  168047
168048
0      0.244928  0.450782  0.002506  0.031404  ...    0.0    0.0    0.0
0.0
1      0.330435  0.558513  0.003006  0.022795  ...    0.0    0.0    0.0
0.0
2      0.159420  0.302915  0.005137  0.018988  ...    0.0    0.0    0.0
0.0
3      0.136232  0.051542  0.007769  0.018394  ...    0.0    0.0    0.0
0.0
4      0.294928  0.662019  0.002999  0.022111  ...    0.0    0.0    1.0
0.0
...          ...          ...          ...          ...          ...          ...
...
```

```

487915  0.275362  0.028306  0.004646  0.007792  ...    0.0    0.0    0.0
0.0
487916  0.217391  0.239544  0.009456  0.035657  ...    0.0    0.0    0.0
0.0
487917  0.610145  0.434728  0.010822  0.010617  ...    0.0    0.0    0.0
0.0
487918  0.420290  0.324039  0.008296  0.004238  ...    0.0    0.0    0.0
0.0
487919  0.275362  0.260668  0.004611  0.015895  ...    0.0    0.0    0.0
0.0

```

[487920 rows x 168049 columns]

#- Note i wrote a logic tried multiple times to run below nural network but because i was working on colab and its provide only 12gb ram it was got creashed during training so did not perform full training but i kept logic as it is

```

[ ]: X_input = Input(shape=(X_train_dence.shape[1]), name= 'X_input', sparse=True)
Dence0 = Dense(32,activation="relu",kernel_initializer='he_uniform')(X_input)

Dence1 = Dense(32,activation="relu",kernel_initializer='he_uniform')(Dence0)
Dropout1 = Dropout(0.1)(Dence1)
Dence2 = Dense(64,activation="relu",kernel_initializer='he_uniform')(Dropout1)
BnLayer = BatchNormalization()(Dence2)
Dropout2 = Dropout(0.1)(BnLayer)
Dence3 = Dense(64,activation="relu",kernel_initializer='he_uniform')(Dropout2)

output = Dense(1,activation="sigmoid",kernel_initializer='he_uniform')(Dence3)
model = Model(inputs=X_input, outputs=output)

earlystop = EarlyStopping(monitor='aucroc', min_delta=0.0000001, patience=4,
    ↳verbose=2,mode='auto')

optimizer = tf.keras.optimizers.Adam(learning_rate=0.001,beta_1=0.9,beta_2=0.
    ↳999)

model.compile(loss='binary_crossentropy', optimizer=optimizer,
    ↳metrics=['accuracy',aucroc],run_eagerly=True)
save_model = model.fit(
    ↳x=X_train_dence,y=y_train,batch_size=100,epochs=25,validation_data=(X_test_dence,y_test),ca

```

#Load and Predict Funtions

```

[ ]: train_data_preprocessed = pd.read_pickle("train_data_preprocessed.pkl")
train_data_preprocessed.columns

```

```
[ ]: Index(['loan_amnt', 'term', 'batch_enrolled', 'int_rate', 'grade', 'sub_grade',
          'emp_title', 'emp_length', 'home_ownership', 'annual_inc',
          'verification_status', 'purpose', 'title', 'zip_code', 'addr_state',
          'dti', 'delinq_2yrs', 'inq_last_6mths', 'open_acc', 'pub_rec',
          'revol_bal', 'revol_util', 'total_acc', 'initial_list_status',
          'total_rec_int', 'total_rec_late_fee', 'recoveries',
          'collection_recovery_fee', 'collections_12_mths_ex_med',
          'last_week_pay', 'acc_now_delinq', 'tot_coll_amt', 'tot_cur_bal',
          'total_rev_hi_lim', 'loan_status', 'paybility'],
          dtype='object')
```

6 Result summary in tabular form

Model

Hyperparameter

Validation Auc

Logistic regression classifier

$C = 1$

0.879

GBDT classifier

$n_estimators = 200$, $learning_rate = 0.3$

0.9398

Random Forest

$n_estimators = 100$, $max_depth = 50$

0.839

6.0.1 so gbdt shows best auc score of 0.9398 with least overfitting so we will choose gbdt model to predict our submissions

```
[ ]: def text_operations(features):
      for i in features:
          train_data_preprocessed[i] = train_data_preprocessed[i].map(lambda x: x.
          ↪lower().replace(' ', '_'))

[ ]: text_operations(['term', 'batch_enrolled', 'grade', 'sub_grade', 'emp_title', 'emp_length',
                    'home_ownership', 'verification_status', 'purpose', 'title',
                    'zip_code', 'addr_state', 'initial_list_status',
                    'last_week_pay'])
```

```
[ ]: def predict_payability(row):
    annual_inc = row.annual_inc
    if row.annual_inc == 0:
        #small value to ignore log(0) error
        annual_inc = 0.0001
    return (math.log(row.loan_amnt + (float(row.term.split('_')[0]) * row.
    ↪int_rate)) + 0.00001 )/math.log(annual_inc)

[ ]: import math
from scipy.sparse import hstack
from sklearn.externals import joblib
def load_predict():
    test_data = pd.read_csv('test_indessa.csv')
    member_id = test_data.member_id

    # select selected features which is in train_preprocessing
    test_data = test_data[['loan_amnt', 'term', 'batch_enrolled', 'int_rate',
    ↪'grade', 'sub_grade',
        'emp_title', 'emp_length', 'home_ownership', 'annual_inc',
        'verification_status', 'purpose', 'title', 'zip_code', 'addr_state',
        'dti', 'delinq_2yrs', 'inq_last_6mths', 'open_acc', 'pub_rec',
        'revol_bal', 'revol_util', 'total_acc', 'initial_list_status',
        'total_rec_int', 'total_rec_late_fee', 'recoveries',
        'collection_recovery_fee', 'collections_12_mths_ex_med',
        'last_week_pay', 'acc_now_delinq', 'tot_coll_amt', 'tot_cur_bal',
        'total_rev_hi_lim']]

    #replace nan and black spaces in batch enrolled
    test_data['batch_enrolled'] = test_data['batch_enrolled'].fillna('unknown')
    test_data['batch_enrolled'] = test_data['batch_enrolled'].map(lambda x:
    ↪'blank' if x == ' ' else x)

    # fill nan values
    median_value_features = pickle.load(open('save_feature_median_value.pickle',
    ↪'rb'))
    max_value_features = pickle.load(open('save_feature_max_value.pickle', 'rb'))
    for i in median_value_features.keys():
        test_data[i] = test_data[i].fillna(median_value_features[i])
    for j in max_value_features.keys():
        test_data[j] = test_data[j].fillna(max_value_features[j])

    #preprocess categorica data
    cat_features =
    ↪['term', 'batch_enrolled', 'grade', 'sub_grade', 'emp_title', 'emp_length'
        , 'home_ownership', 'verification_status', 'purpose', 'title'
        , 'zip_code', 'addr_state', 'initial_list_status'
        , 'last_week_pay']
```

```

for i in cat_features:
    test_data[i] = test_data[i].map(lambda x: x.lower().replace(' ', '_'))

#create new feature
test_data['paybility'] = test_data.apply(lambda row: predict_paybility(row),
→axis=1)

#normalize numerical features
numerical_features = ['loan_amnt', 'int_rate',
    'annual_inc', 'dti', 'delinq_2yrs', 'inq_last_6mths', 'open_acc',
    'pub_rec', 'revol_bal', 'revol_util', 'total_acc', 'total_rec_int',
    'total_rec_late_fee', 'recoveries', 'collection_recovery_fee',
    'collections_12_mths_ex_med', 'acc_now_delinq', 'tot_coll_amt',
    'tot_cur_bal', 'total_rev_hi_lim', 'paybility']
normalizer_scaler = pickle.load(open('normalizers_numeric.pkl', 'rb'))
test_data_numerical_features = test_data[numerical_features].values
test_data_numerical = normalizer_scaler.
→transform(test_data_numerical_features)

#one hot encode categorical features
categorical_features =
→['term', 'batch_enrolled', 'grade', 'sub_grade', 'emp_title', 'emp_length'
    , 'home_ownership', 'verification_status', 'purpose', 'title'
    , 'zip_code', 'addr_state', 'initial_list_status'
    , 'last_week_pay']
for i in categorical_features:
    vector = pickle.load(open(f'vectorizers/vectorizer_{i}.pkl', 'rb'))
    test_feature = vector.transform(test_data[i].values.reshape(-1,1))
    globals()[f"test_feature_{i}"] = test_feature

test_data_final = hstack((test_data_numerical , test_feature_term
→, test_feature_batch_enrolled, test_feature_grade, test_feature_sub_grade, test_feature_emp_tit
    ,
→, test_feature_home_ownership, test_feature_verification_status, test_feature_purpose, test_fea
    ,
→, test_feature_zip_code, test_feature_addr_state, test_feature_initial_list_status
    , test_feature_last_week_pay )).tocsr()

#predict probabilities
xgbClf = joblib.load('xgb_model.joblib')
y_probabilities = batch_predict(xgbClf , test_data_final )
test_submission_df = pd.DataFrame({'member_id':member_id, 'loan_status':
→y_probabilities})
test_submission_df.to_csv('test_submission.csv')
test_submission_df.head()

```

```
load_predict()
```

```
[ ]:  member_id  loan_status
0    11937648    0.410839
1    38983318    0.291551
2    27999917    0.164921
3    61514932    0.206364
4    59622821    0.304588
```

```
[ ]:
```